

Recurrent neural networks for natural language



Roger Levy

9.19: Computational Psycholinguistics

3 November 2021

Agenda for the day

- Last time: with a hidden layer, a NN can learn XOR...



- ...but language isn't just 2D input+2-class output! So, **today**:
- Dealing with language in neural networks
- Recurrent neural networks (RNNs)
 - Simple recurrent networks (SRNs)
 - Gated recurrent units (GRUs)
 - Long short-term memory networks (LSTMs)
- Examining RNN behavior

Dealing with language inputs

Adam adores zebras . . .

Dealing with language inputs

For language, input $\{x_i\}$ and output prediction y seem discrete:

Adam adores zebras . . .

Dealing with language inputs

For language, input $\{x_i\}$ and output prediction y seem discrete:

Adam adores zebras . . .

Simplest approach is ***localist*** or ***one-hot*** representations:

$$\begin{array}{ccc} \text{Adam} \rightarrow & \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} & \text{adores} \rightarrow \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} & \text{zebras} \rightarrow \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 1 \end{bmatrix} \end{array}$$

Dealing with language inputs

For language, input $\{x_i\}$ and output prediction y seem discrete:

Adam adores zebras ...

Simplest approach is **localist** or **one-hot** representations:

$$\begin{array}{ccc} \text{Adam} \rightarrow & \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} & \text{adores} \rightarrow \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} & \text{zebras} \rightarrow \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 1 \end{bmatrix} \end{array}$$

But lower-dimensional **embeddings** capture word similarities:

$$\text{Adam} \rightarrow \begin{bmatrix} 0.6 \\ 0.3 \end{bmatrix} \quad \text{adores} \rightarrow \begin{bmatrix} -0.3 \\ 0.4 \end{bmatrix} \quad \text{zebras} \rightarrow \begin{bmatrix} 0.7 \\ -0.1 \end{bmatrix}$$

Example feed-forward+embedding LM

Bengio et al., 2003: Neural n -gram language model




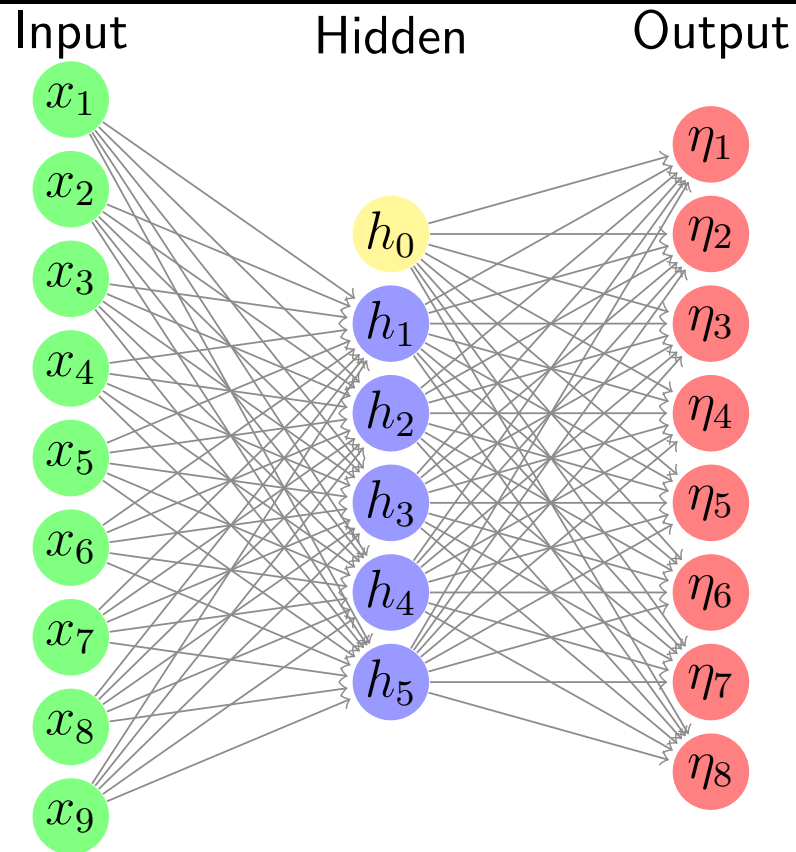
Figure 1: Neural architecture: $f(i, w_{t-1}, \dots, w_{t-n+1}) = g(i, C(w_{t-1}), \dots, C(w_{t-n+1}))$ where g is the neural network and $C(i)$ is the i -th word feature vector.

Old (2003!) perplexity results on Brown corpus

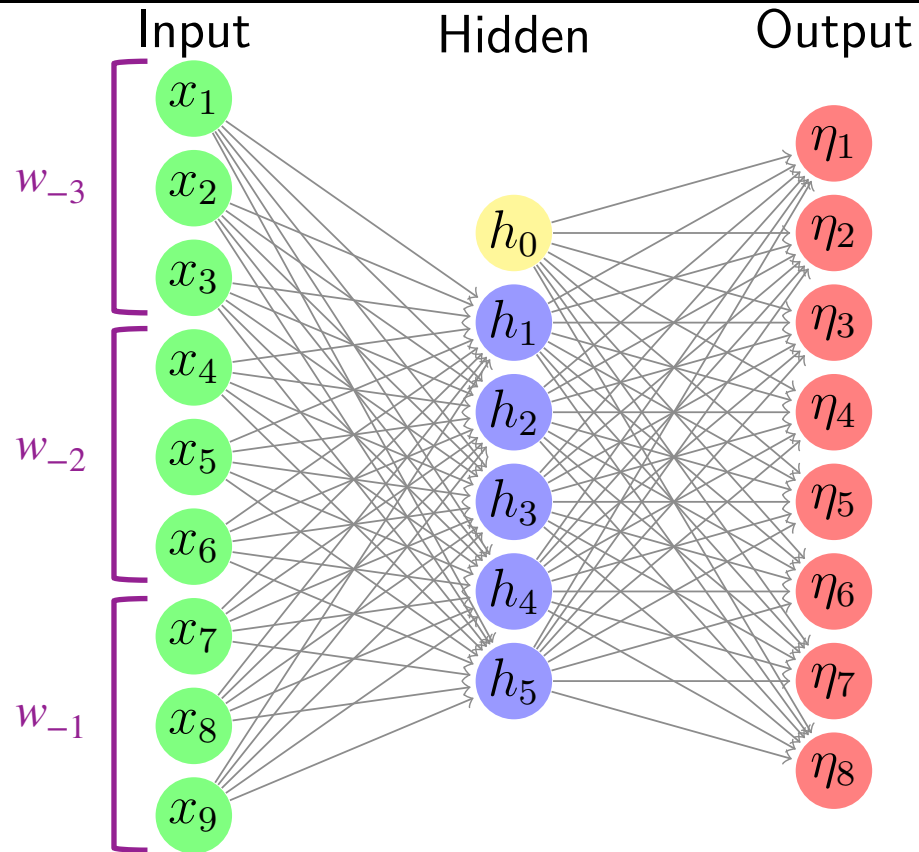
		n	c	h	m	direct	mix	train.	valid.	test.
neural language models	MLP1	5		50	60	yes	no	182	284	268
	MLP2	5		50	60	yes	yes		275	257
	MLP3	5		0	60	yes	no	201	327	310
	MLP4	5		0	60	yes	yes		286	272
	MLP5	5		50	30	yes	no	209	296	279
	MLP6	5		50	30	yes	yes		273	259
	MLP7	3		50	30	yes	no	210	309	293
	MLP8	3		50	30	yes	yes		284	270
	MLP9	5		100	30	no	no	175	280	276
	MLP10	5		100	30	no	yes		265	252
n-gram language models	Del. Int.	3						31	352	336
	Kneser-Ney back-off	3							334	323
	Kneser-Ney back-off	4							332	321
	Kneser-Ney back-off	5							332	321
	class-based back-off	3	150						348	334
	class-based back-off	3	200						354	340
	class-based back-off	3	500						326	312
	class-based back-off	3	1000						335	319
	class-based back-off	3	2000						343	326
	class-based back-off	4	500						327	312
engio et al., 2023)	class-based back-off	5	500						327	312

(Bengio et al.,
2003)

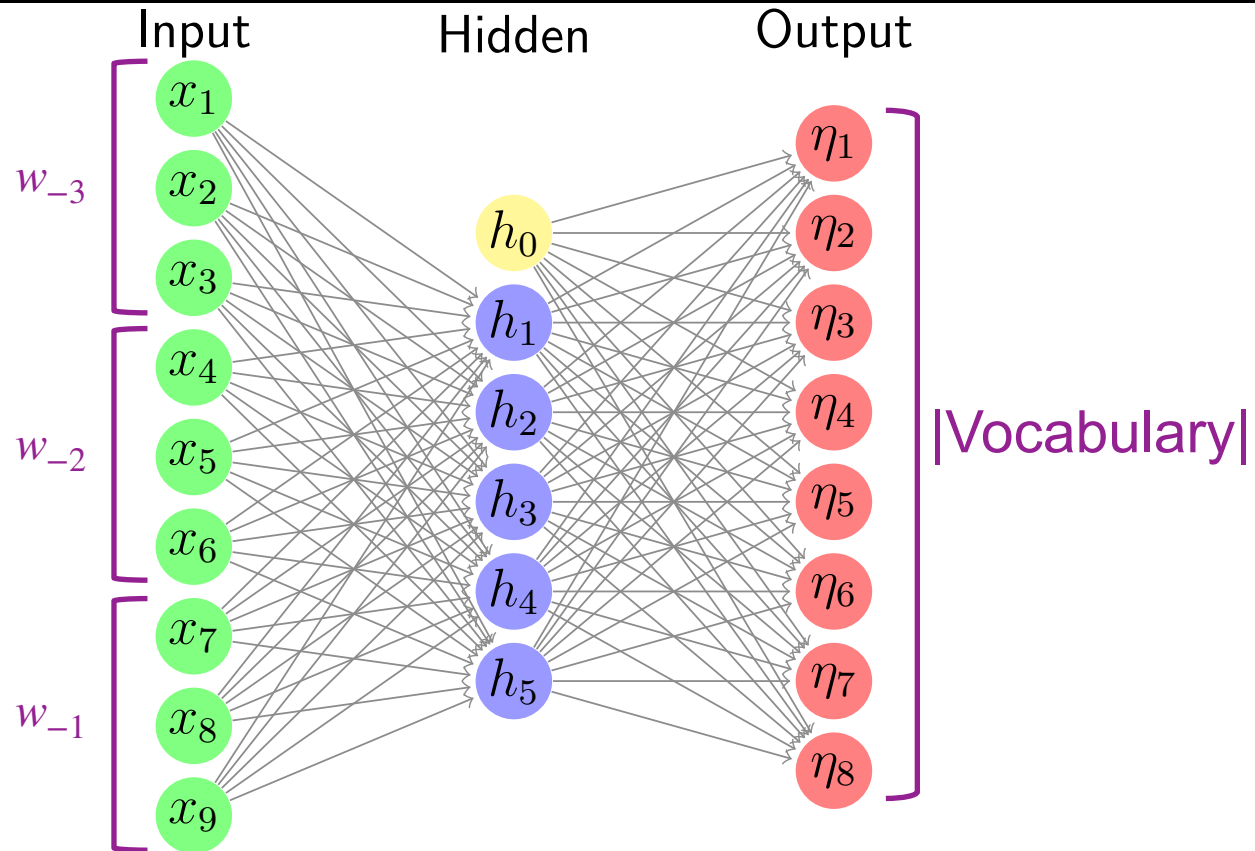
The neural n -gram model




The neural n -gram model



The neural n -gram model

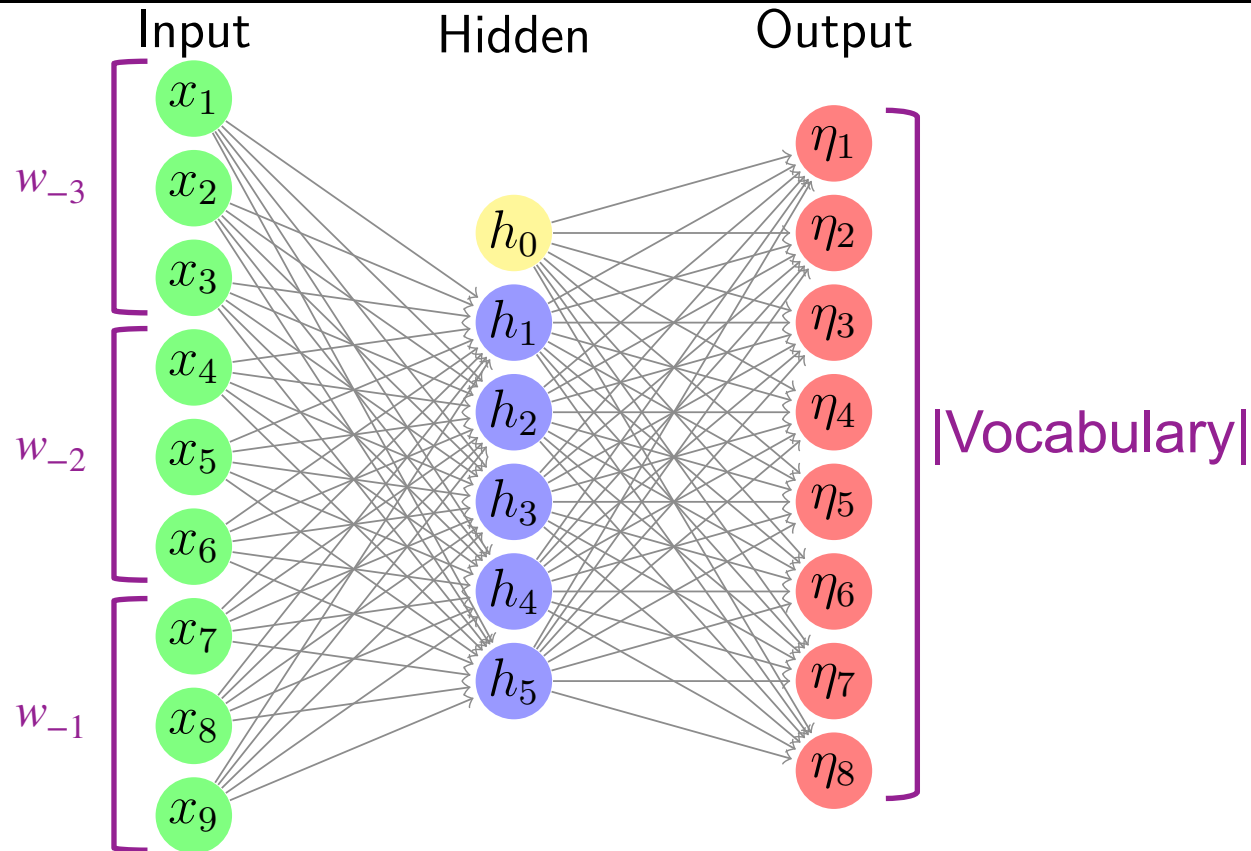


The neural n -gram model



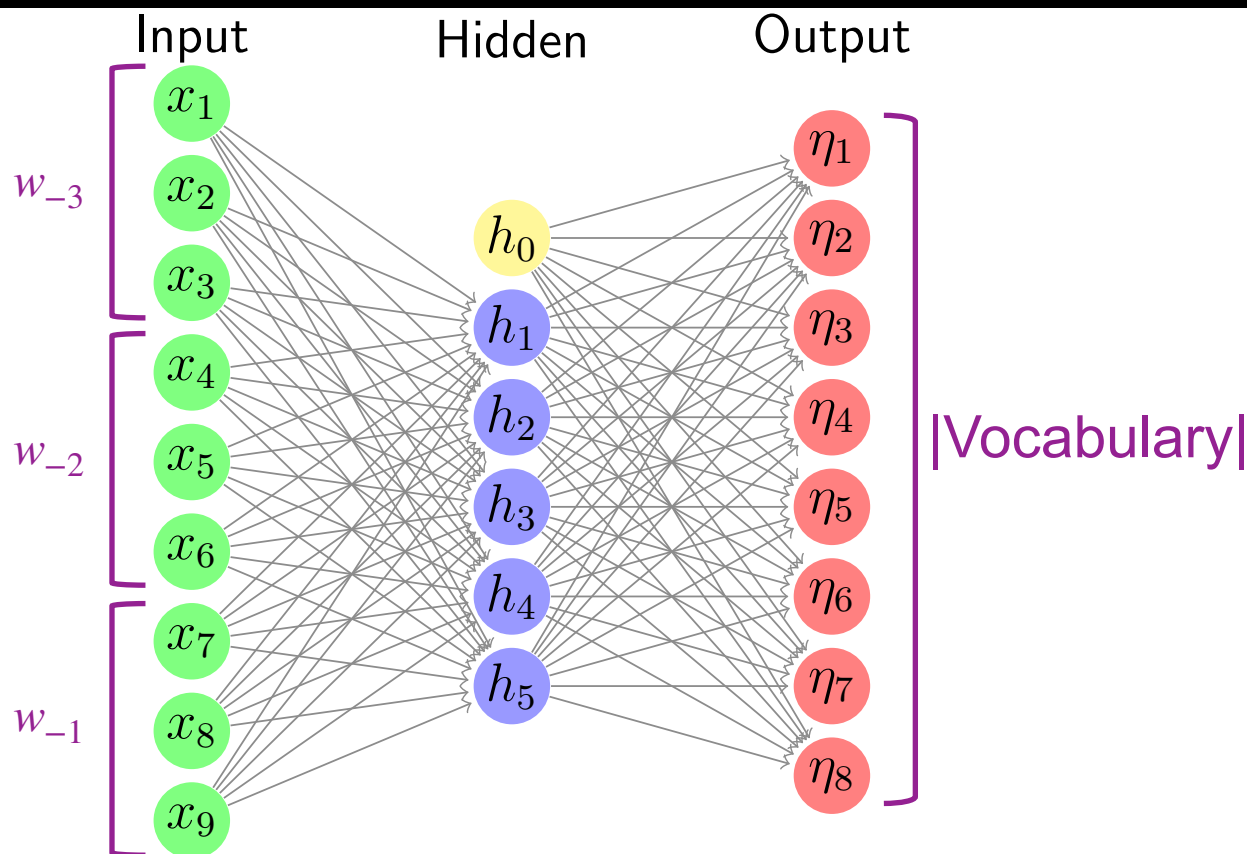
- **Advantages:** generalizes over n -gram contexts

The neural n -gram model



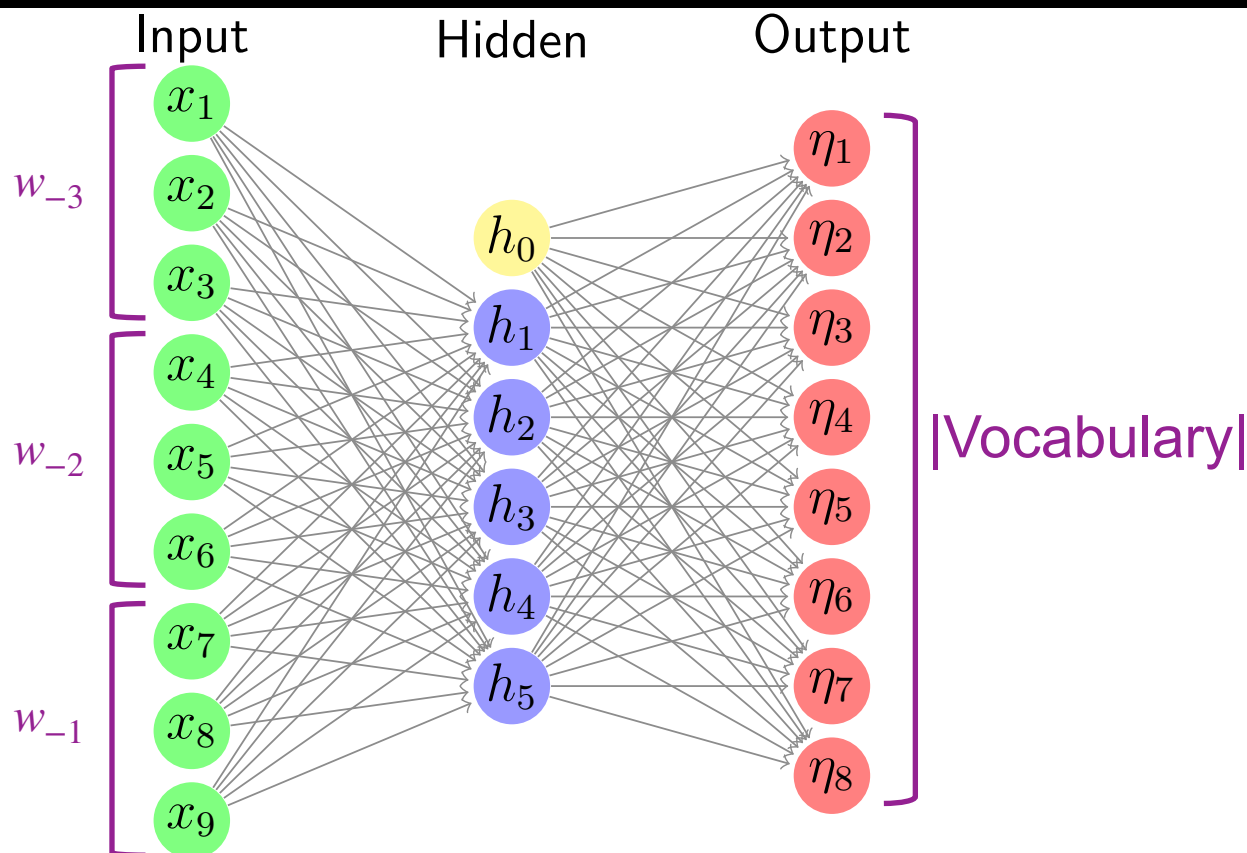
- **Advantages:** generalizes over n -gram contexts
- **Limitations:**

The neural n -gram model



- **Advantages:** generalizes over n -gram contexts
- **Limitations:**
 - this is for a fixed dimensionality input context

The neural n -gram model



- **Advantages:** generalizes over n -gram contexts
- **Limitations:**
 - this is for a fixed dimensionality input context
 - how to model variable-length context, like sentences?

Recurrent neural networks for language

(Jordan, 1986; Elman, 1990)

Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing

Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing
- Previous inputs must be integrated and remembered all together in a uniform representational space


Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing
- Previous inputs must be integrated and remembered all together in a uniform representational space

The woman brought the sandwich from ...

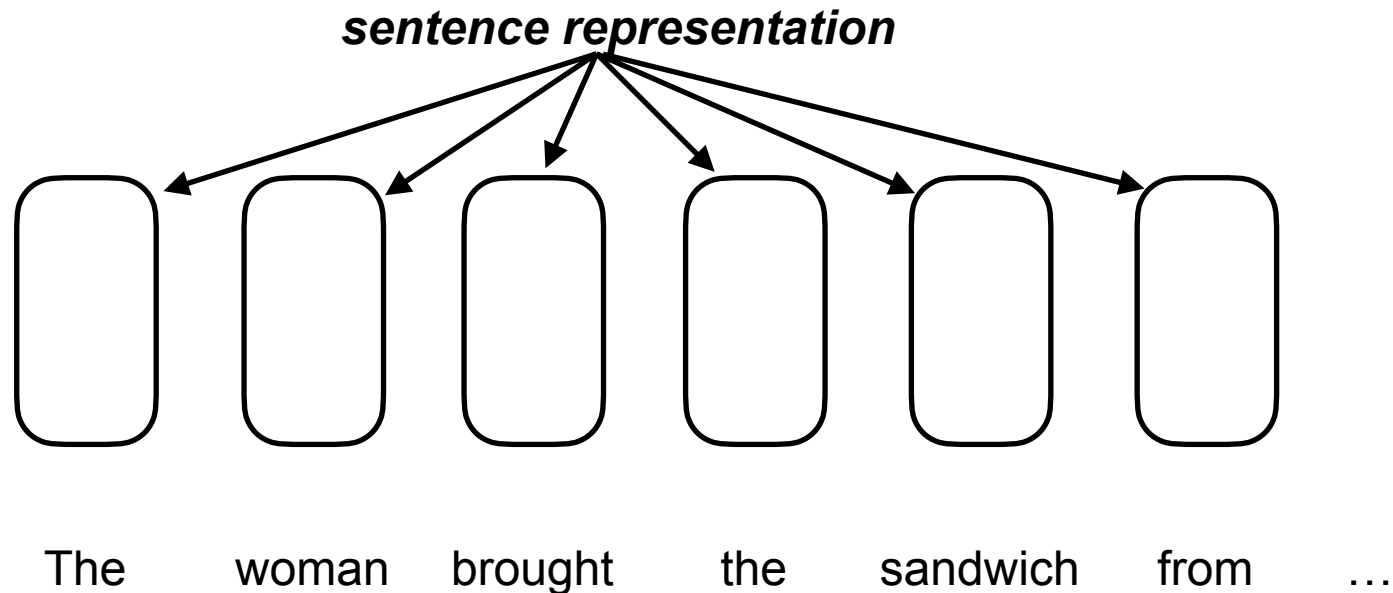
Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing
- Previous inputs must be integrated and remembered all together in a uniform representational space




Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing
- Previous inputs must be integrated and remembered all together in a uniform representational space



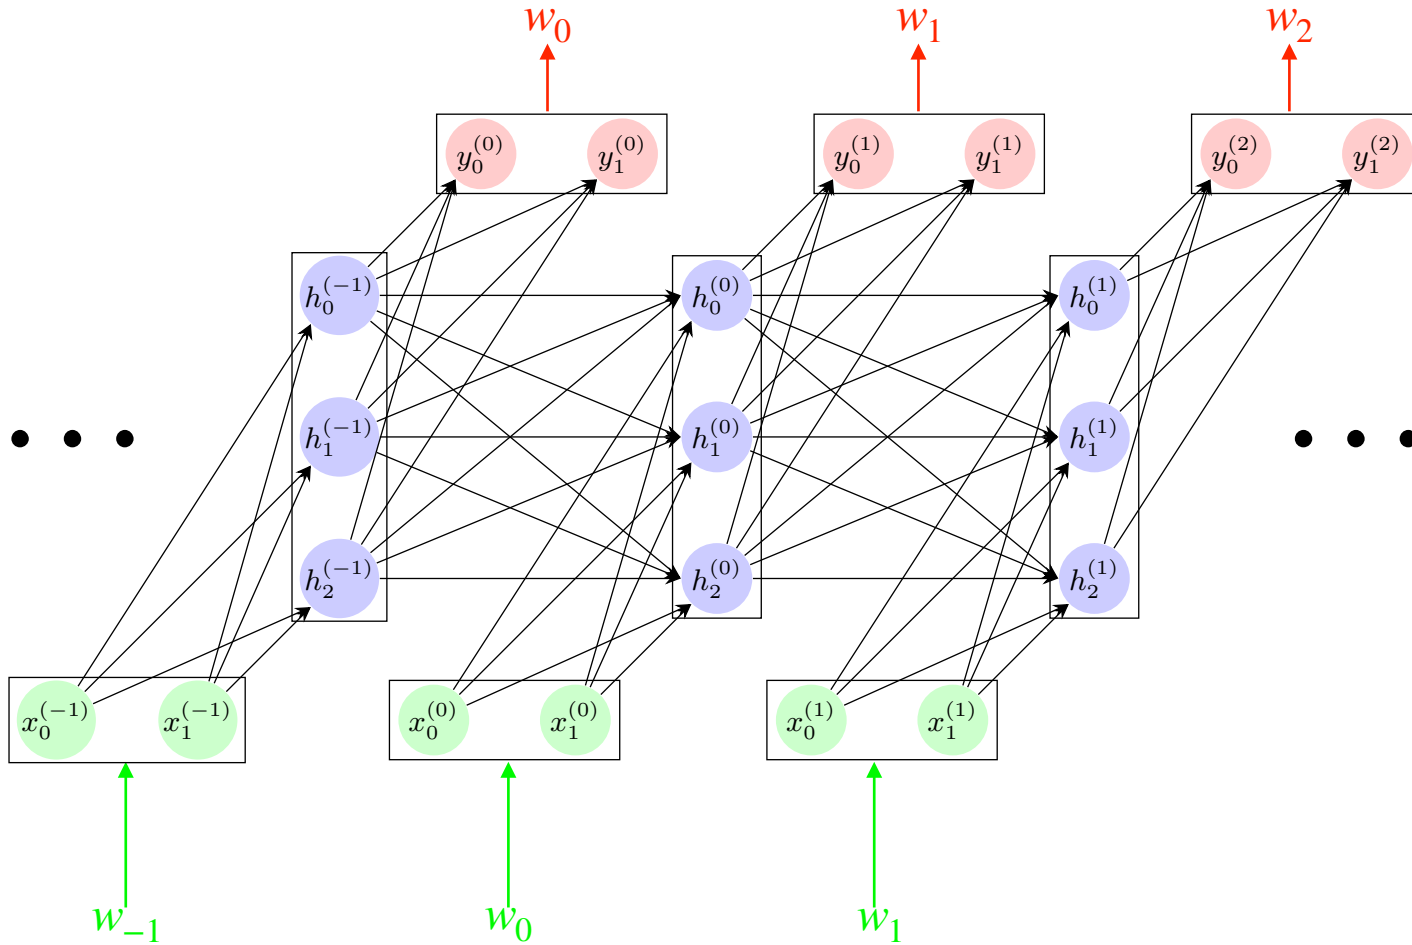
Recurrent neural networks for language

- Draw inspiration from real-time nature of human language processing
- Previous inputs must be integrated and remembered all together in a uniform representational space



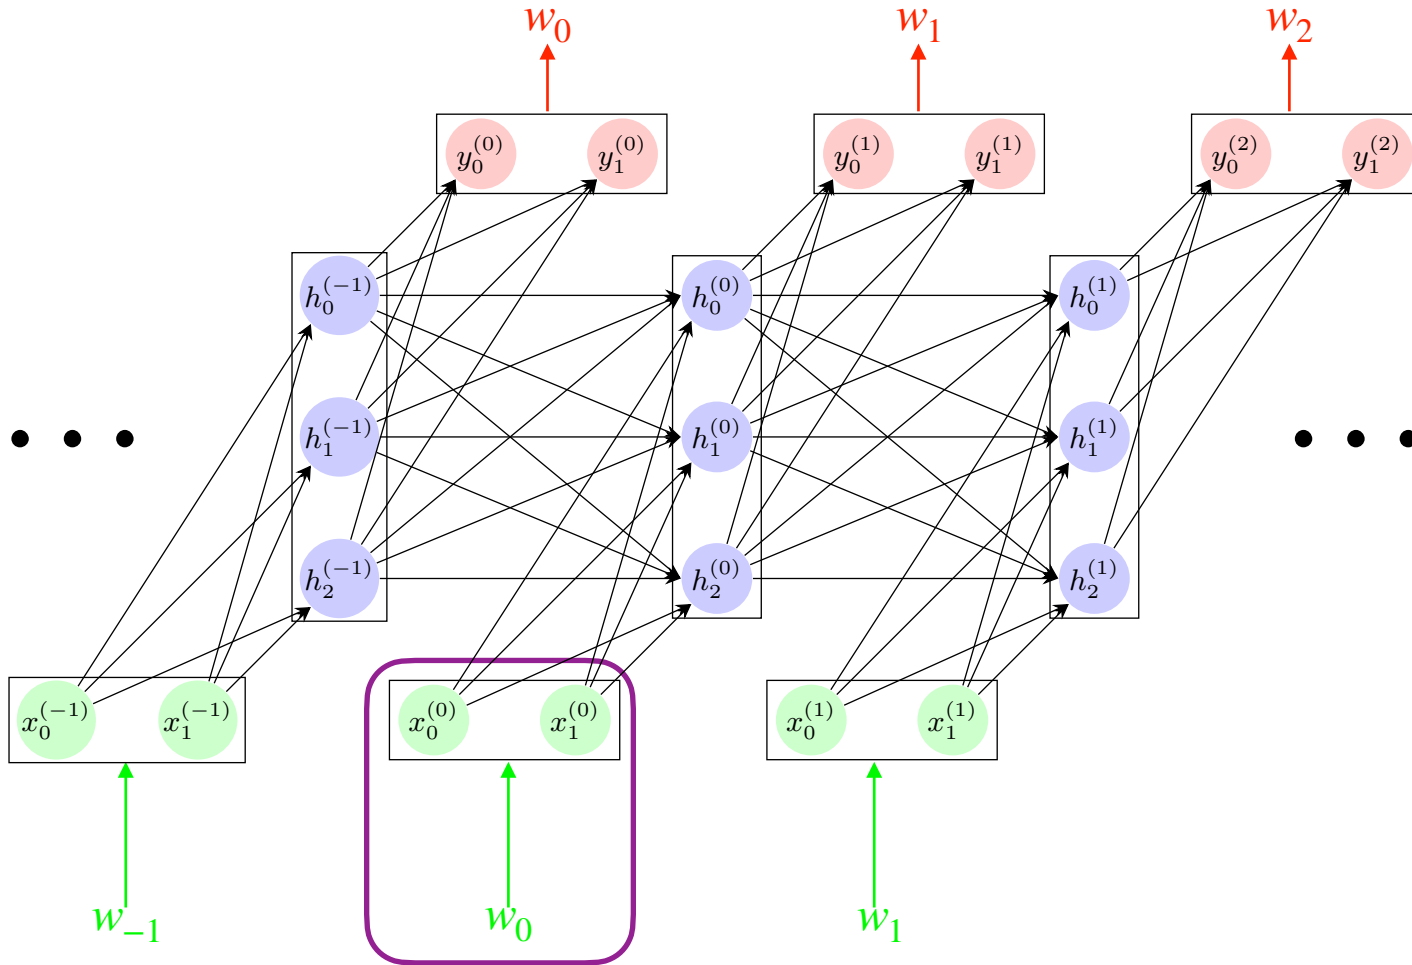
The Simple Recurrent Network (SRN)

(bias nodes not shown)



The Simple Recurrent Network (SRN)

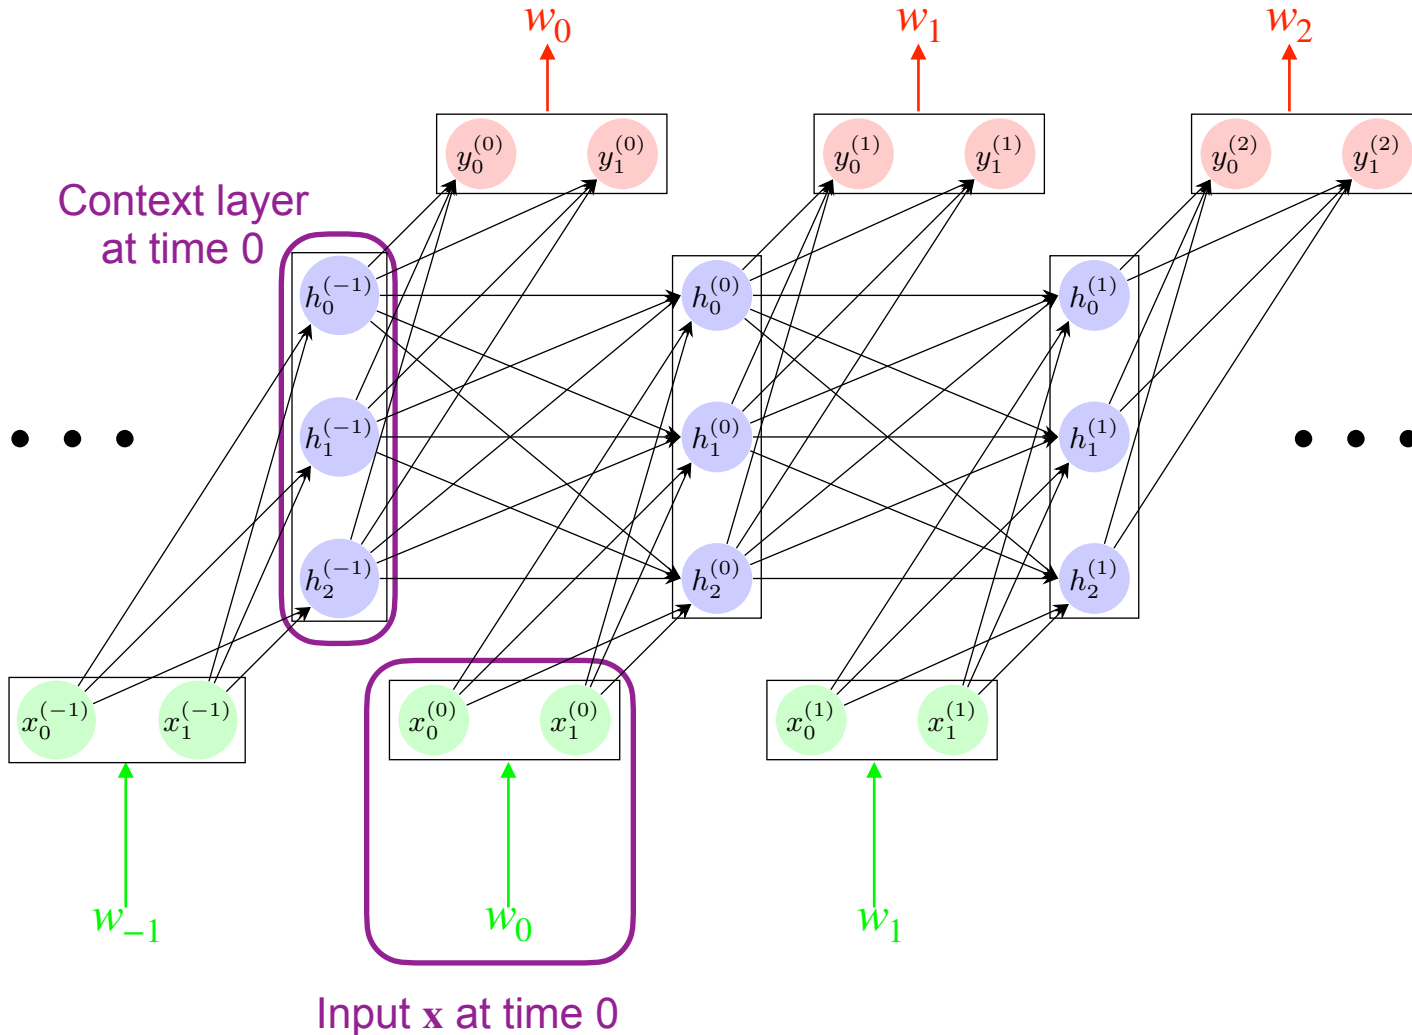
(bias nodes not shown)



Input x at time 0

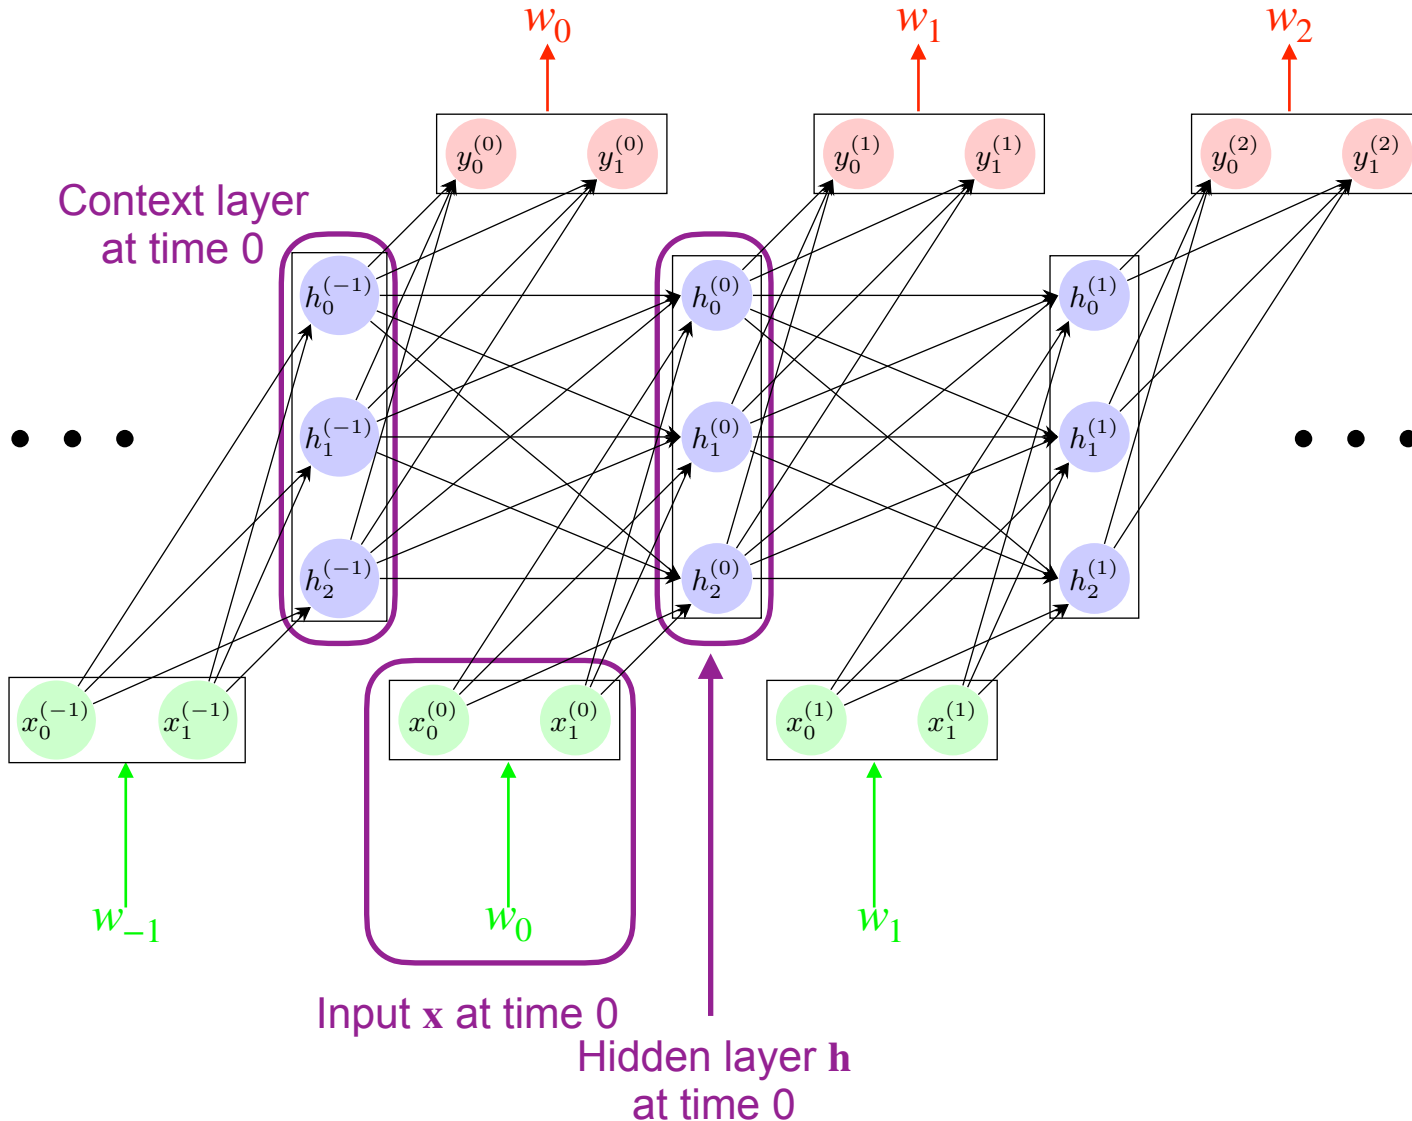
The Simple Recurrent Network (SRN)

(bias nodes not shown)



The Simple Recurrent Network (SRN)

(bias nodes not shown)

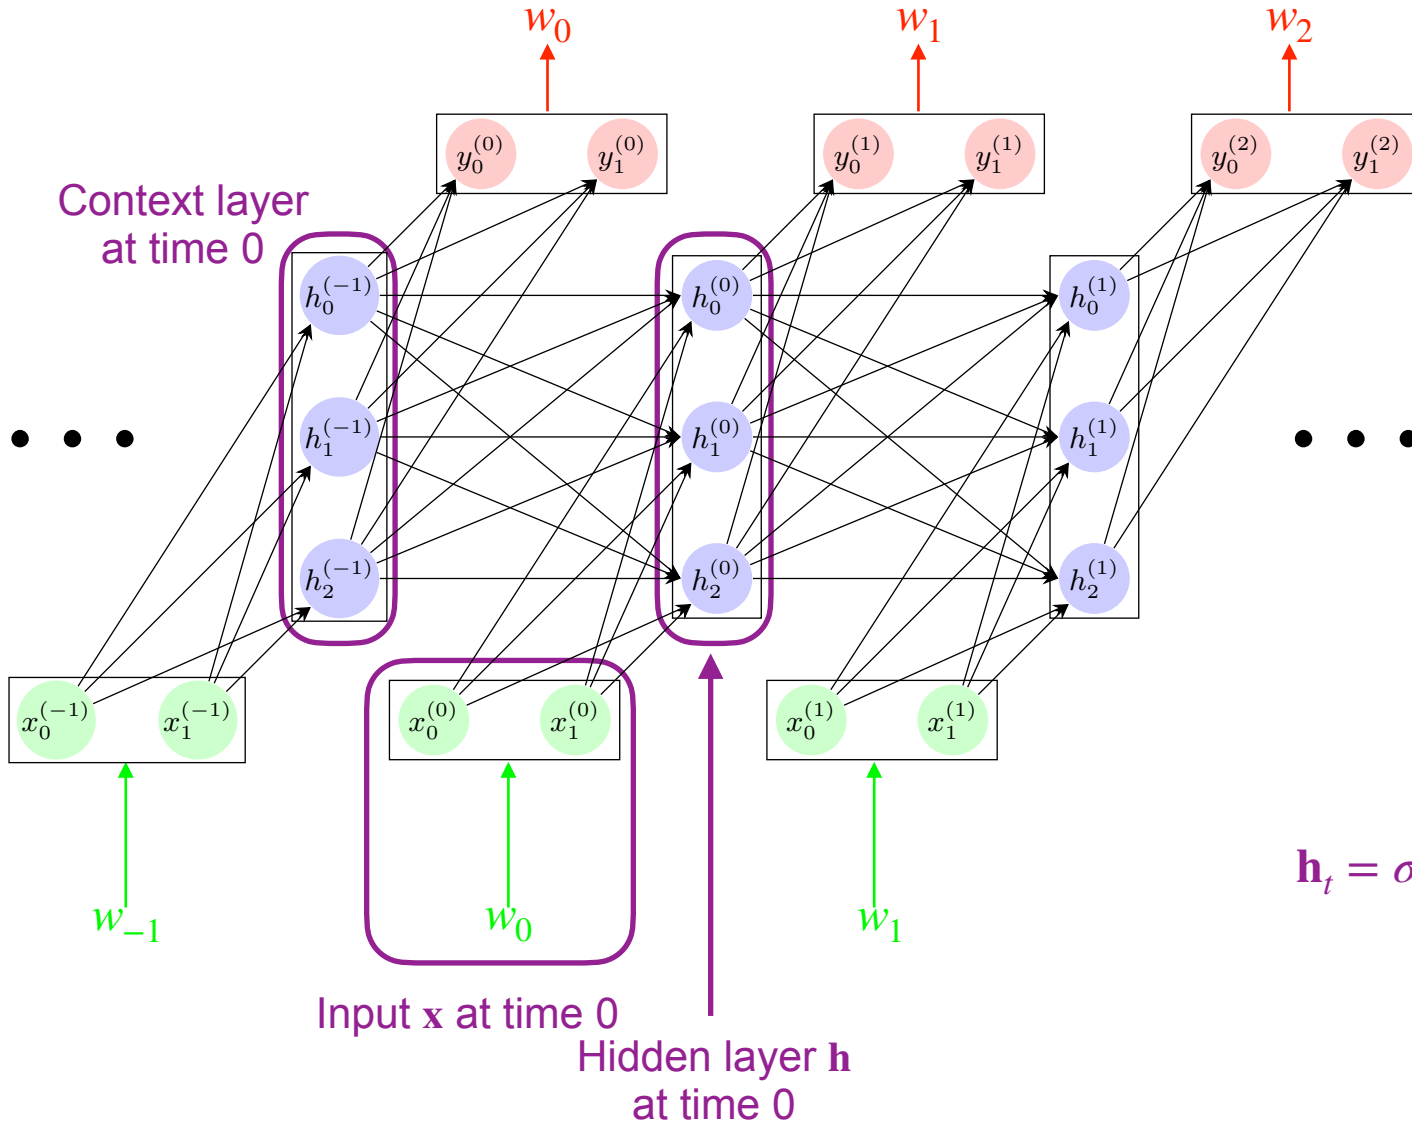


Input x at time 0
Hidden layer h
at time 0

(Elman, 1990)

The Simple Recurrent Network (SRN)

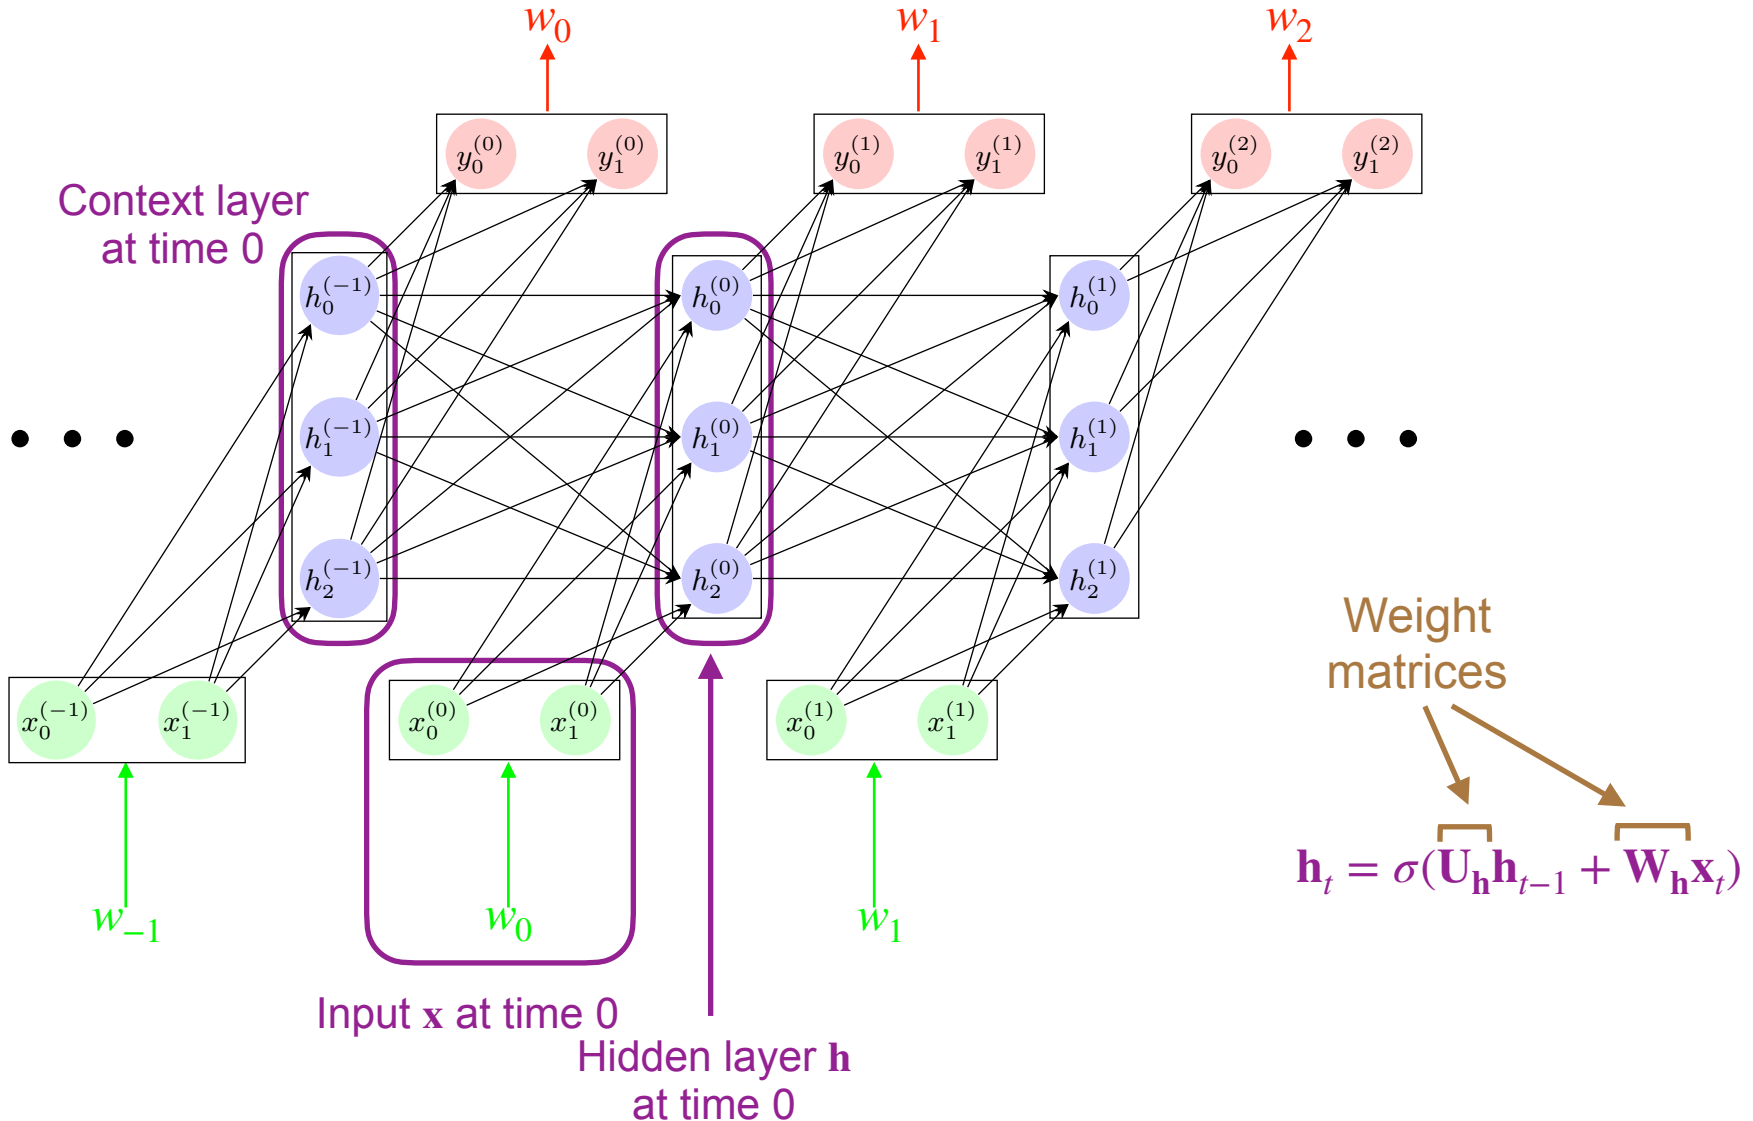
(bias nodes not shown)



$$\mathbf{h}_t = \sigma(\mathbf{U}_h \mathbf{h}_{t-1} + \mathbf{W}_h \mathbf{x}_t)$$

The Simple Recurrent Network (SRN)

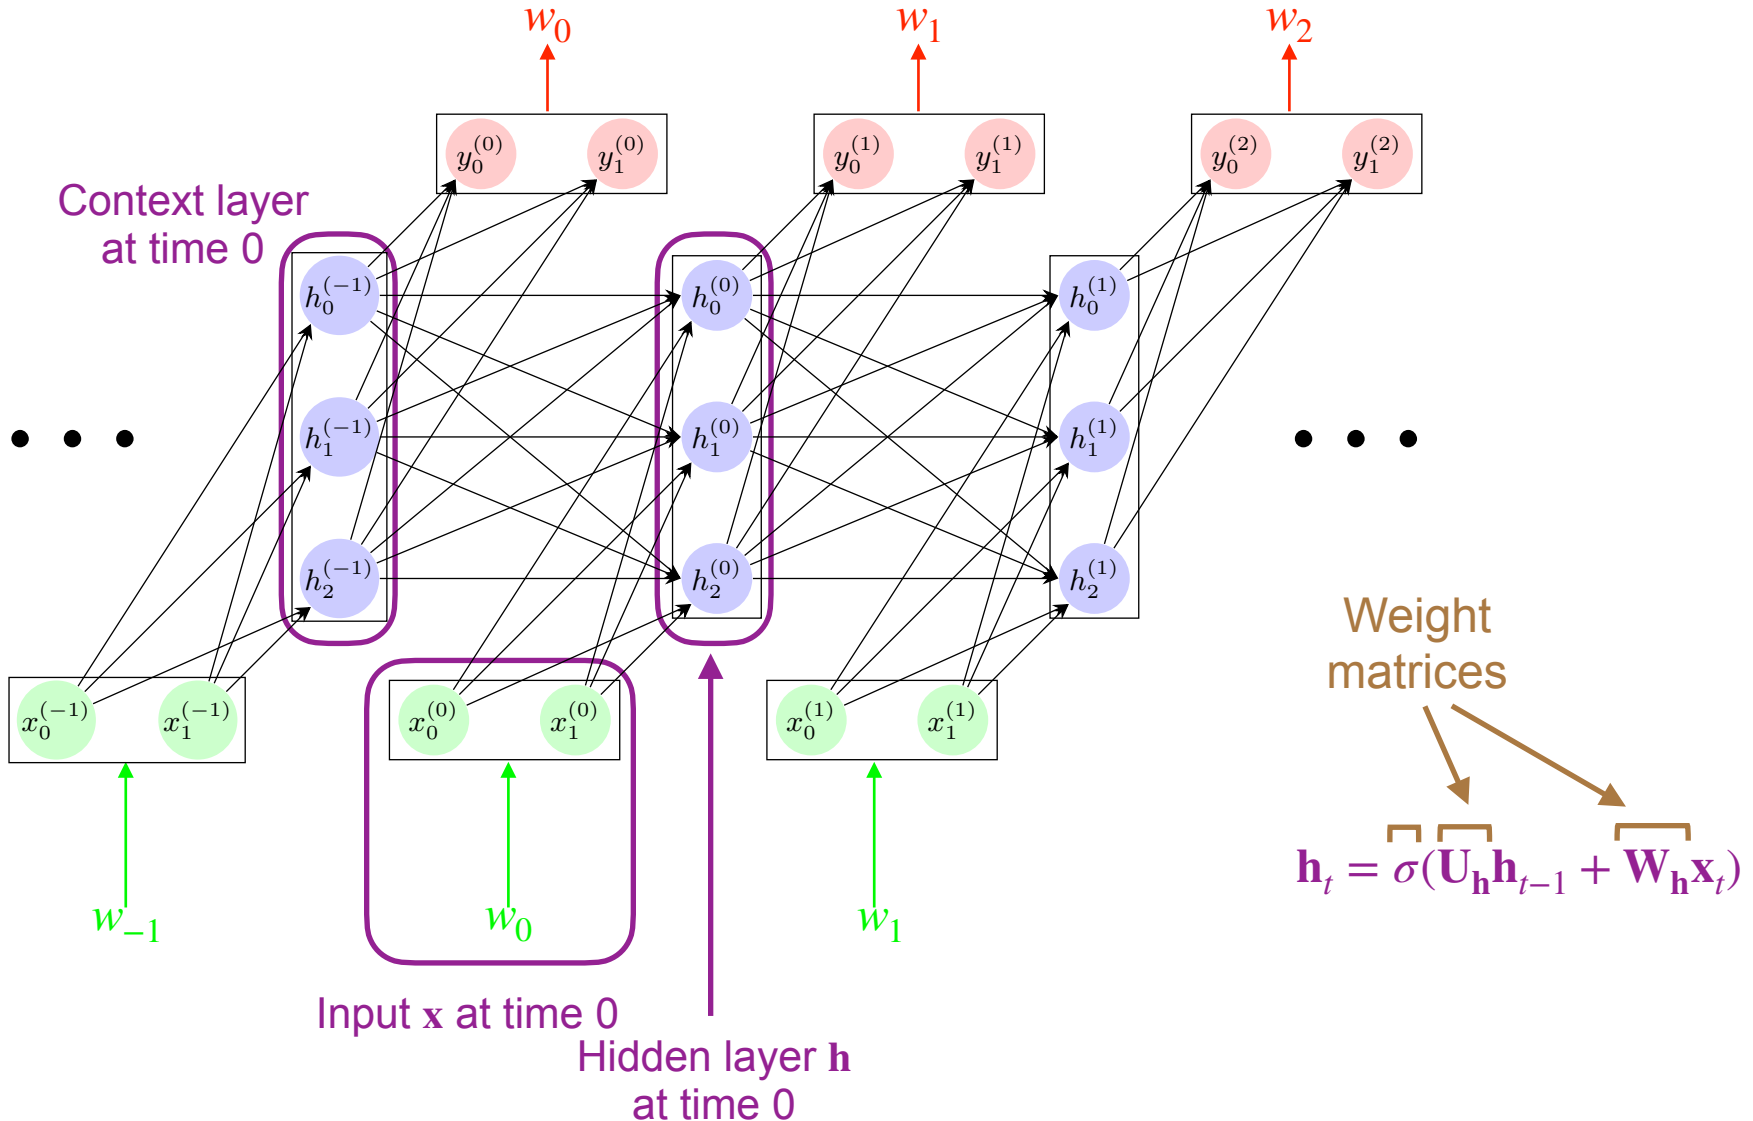
(bias nodes not shown)



(Elman, 1990)

The Simple Recurrent Network (SRN)

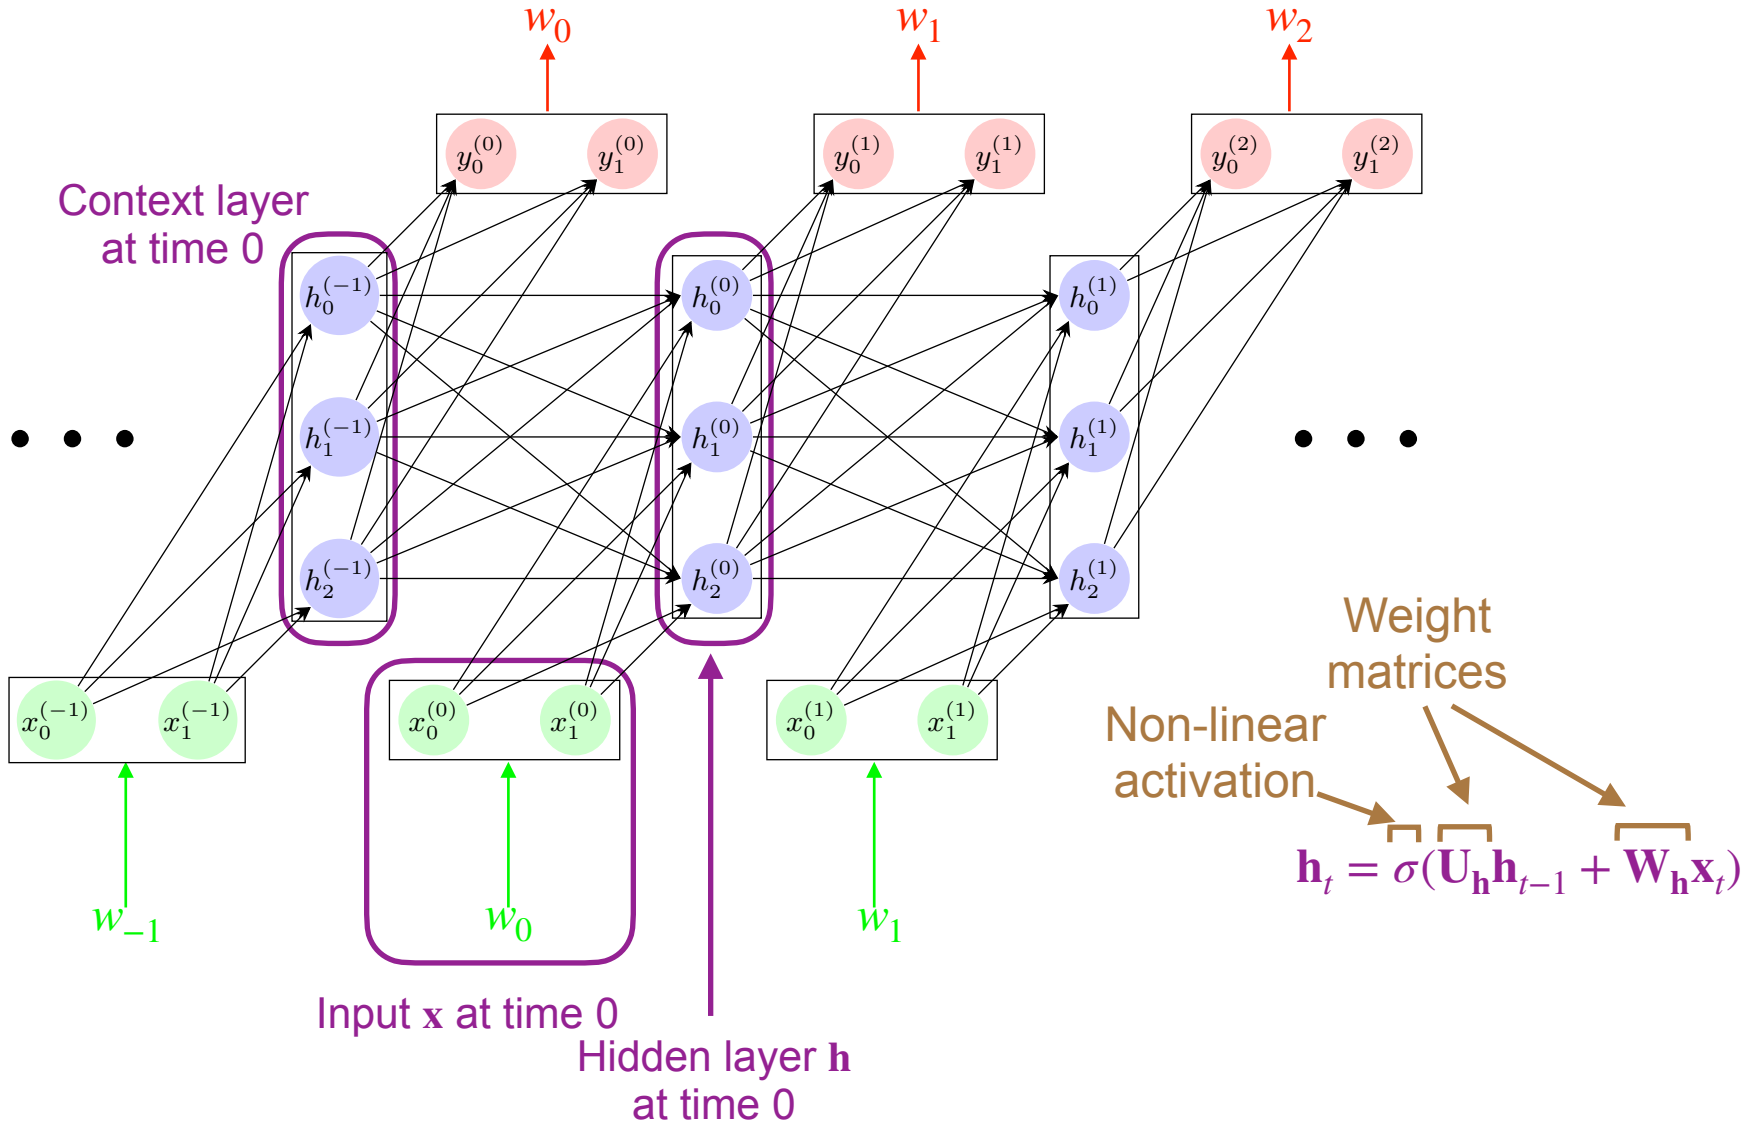
(bias nodes not shown)



(Elman, 1990)

The Simple Recurrent Network (SRN)

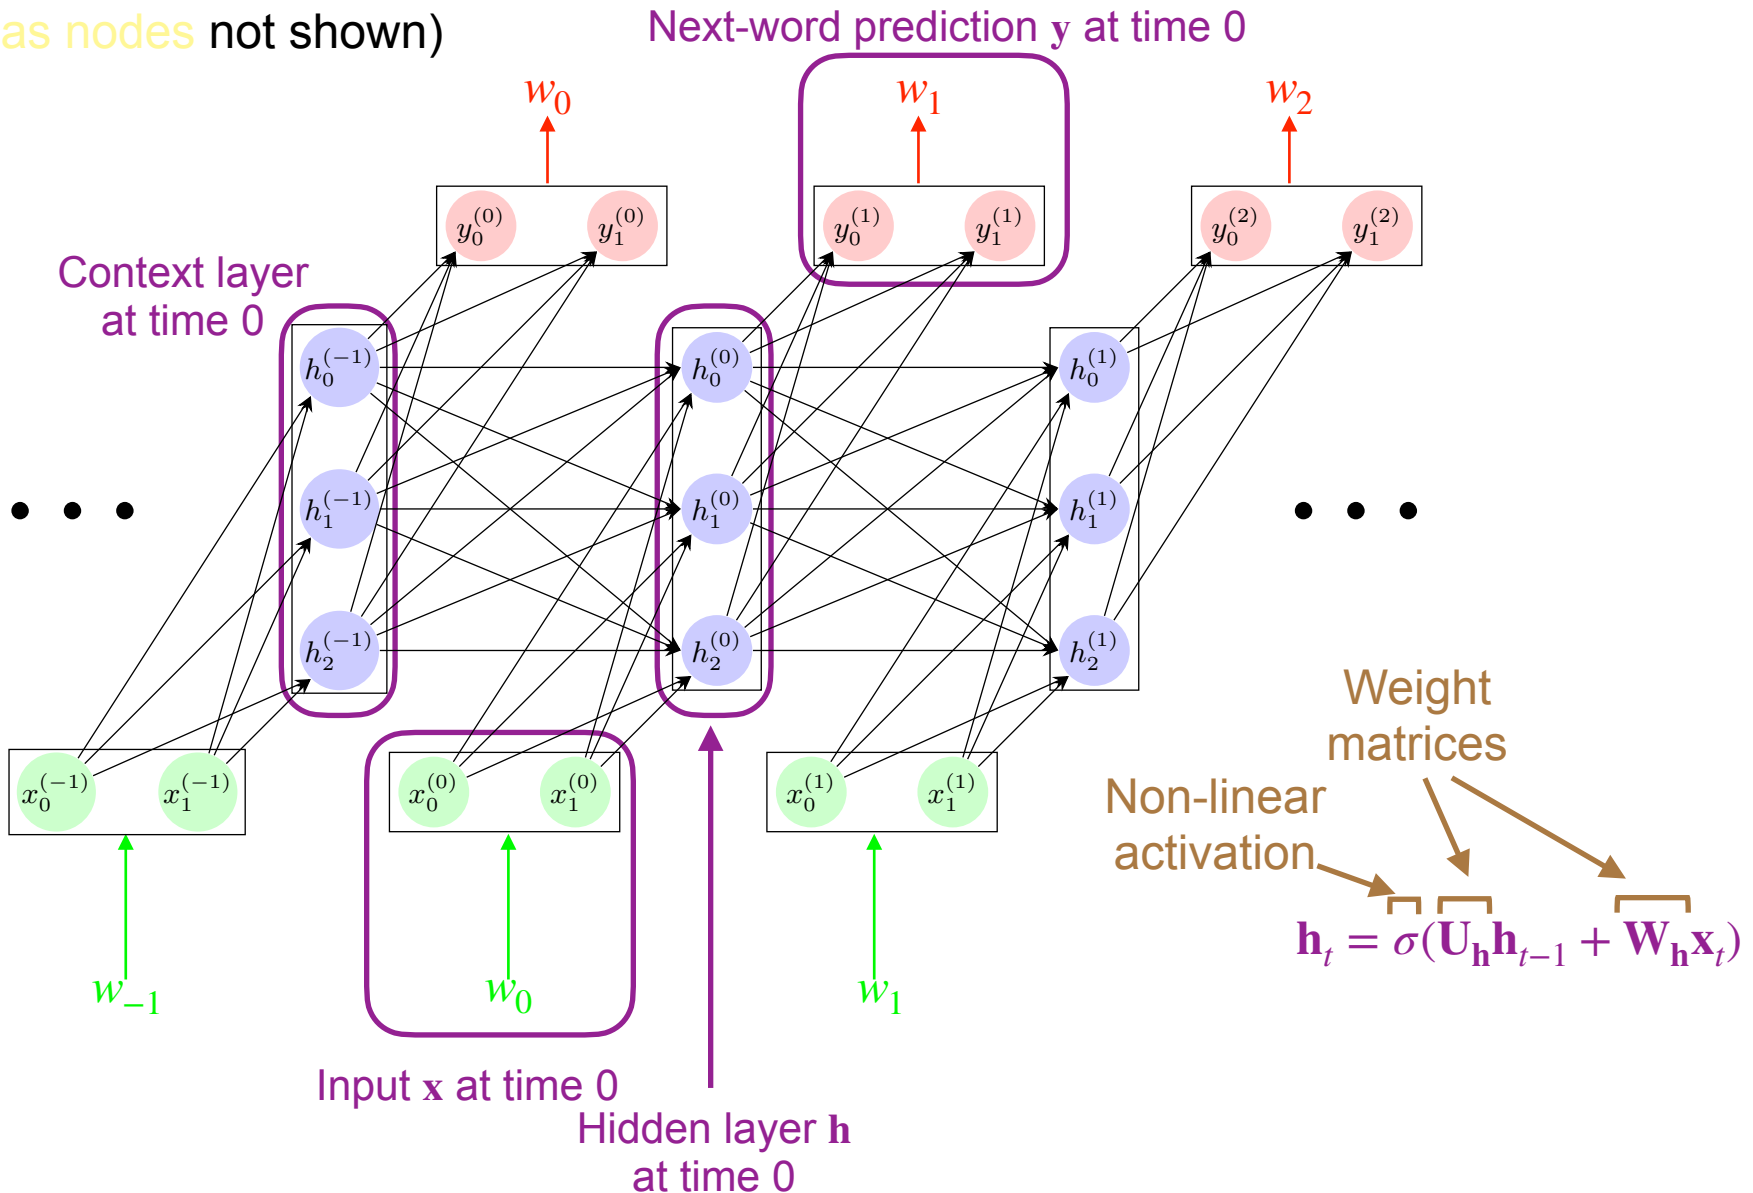
(bias nodes not shown)



(Elman, 1990)

The Simple Recurrent Network (SRN)

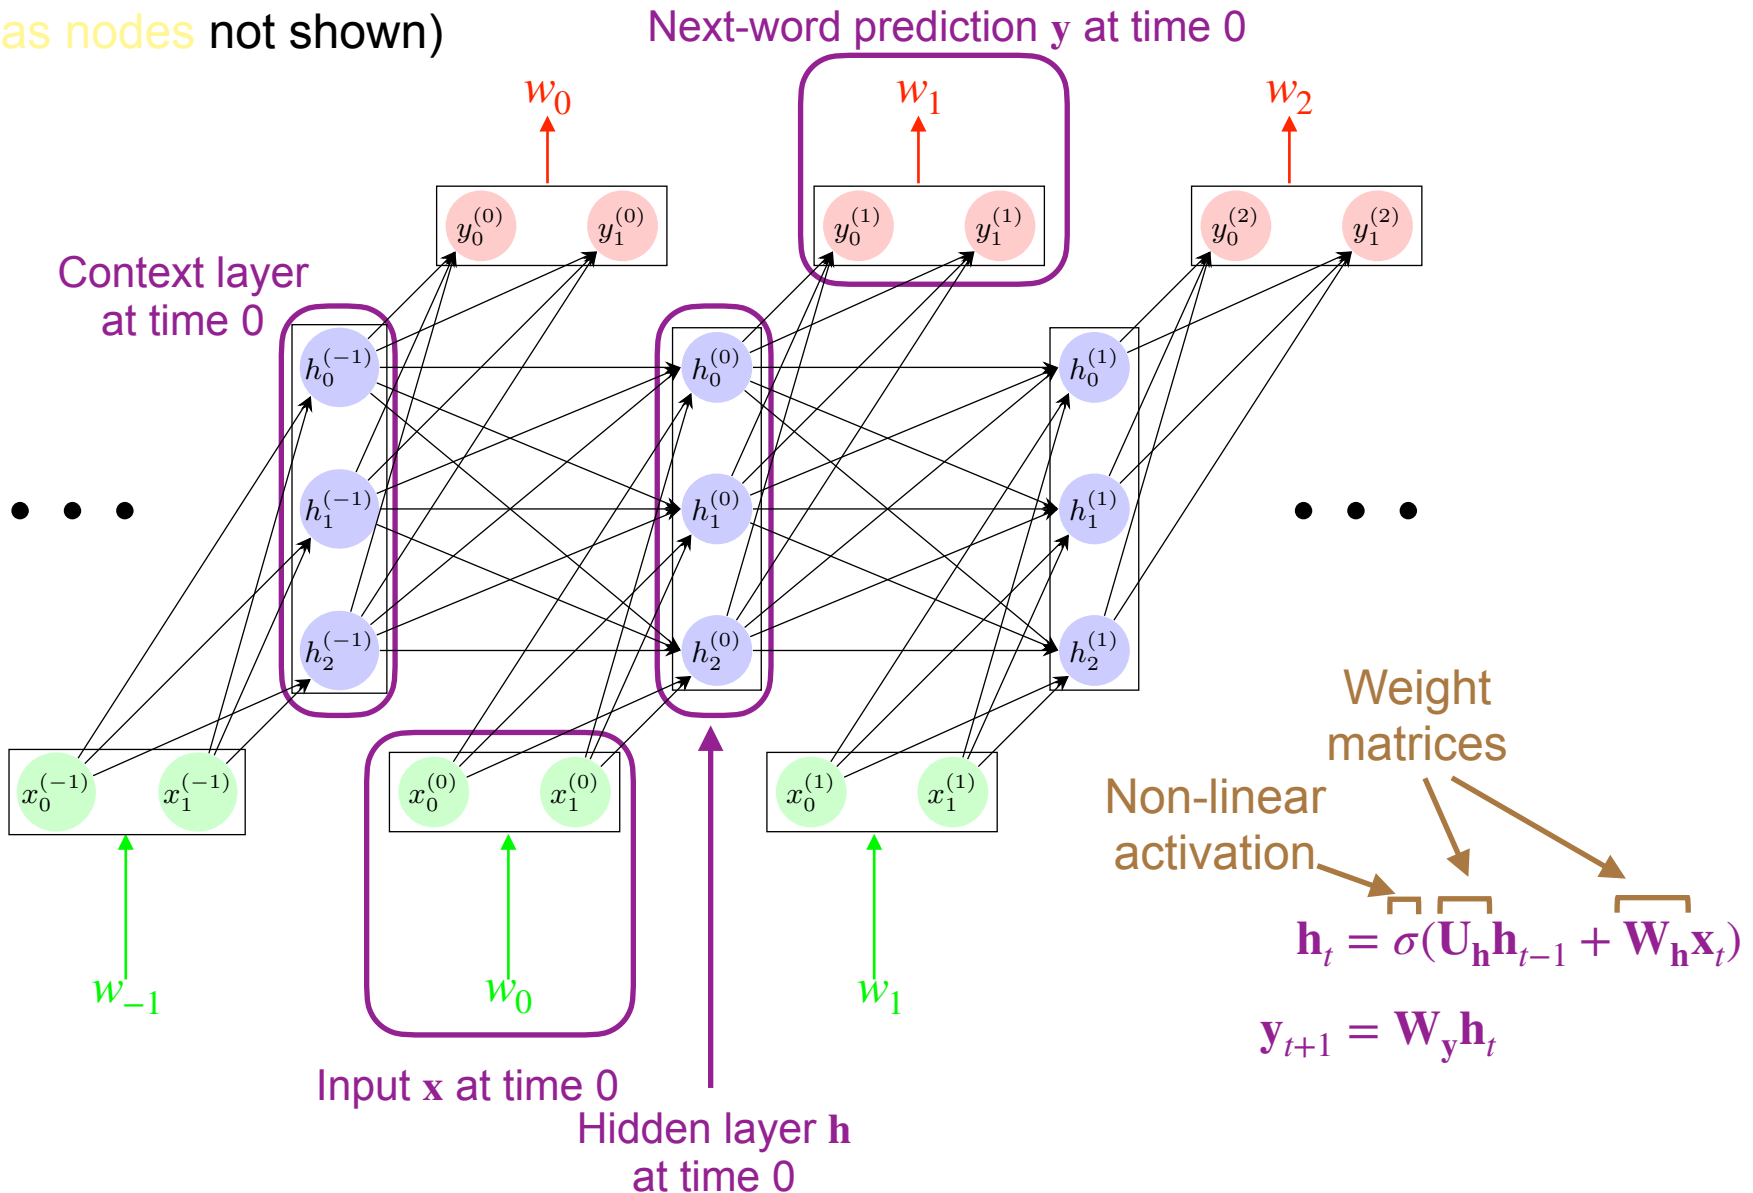
(bias nodes not shown)



(Elman, 1990)

The Simple Recurrent Network (SRN)

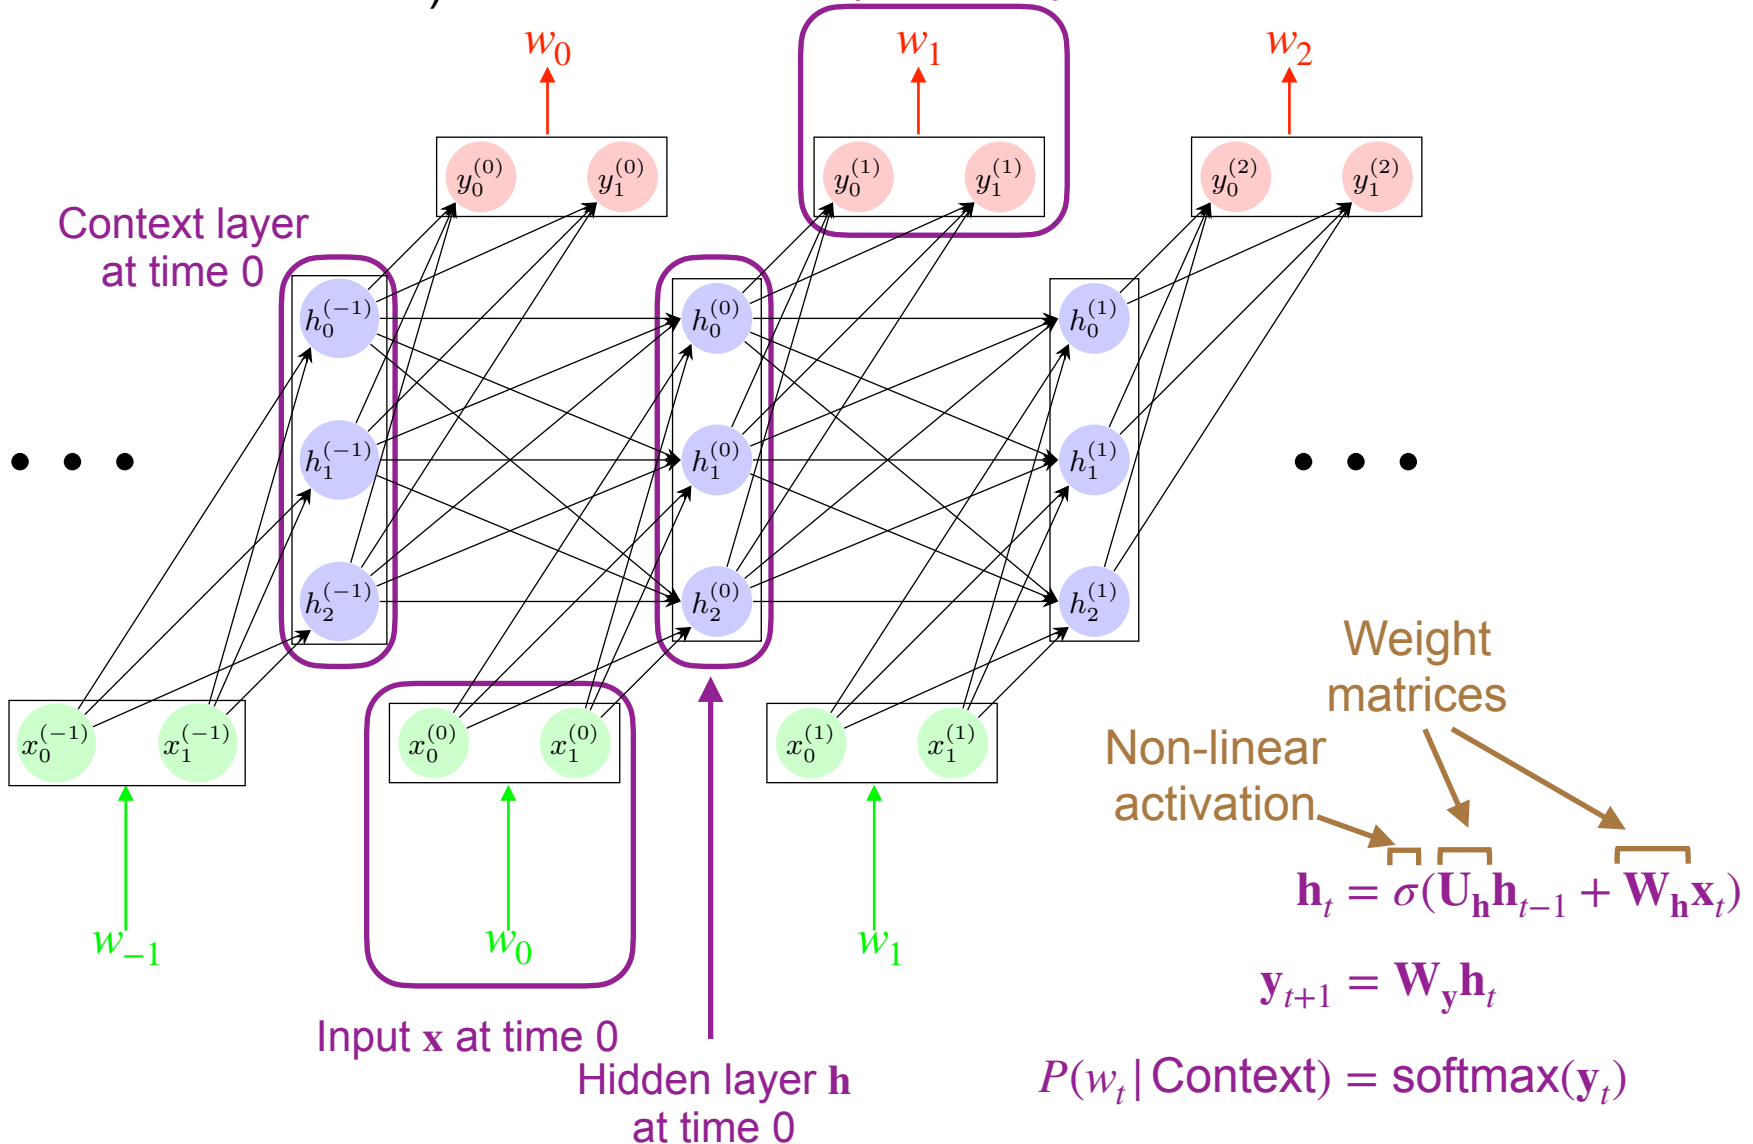
(bias nodes not shown)



The Simple Recurrent Network (SRN)

(bias nodes not shown)


Next-word prediction y at time 0



(Elman, 1990)


SRN "rolled up" and unrolled

- A “rolled-up” representation (Elman, 1990); and unrolled:



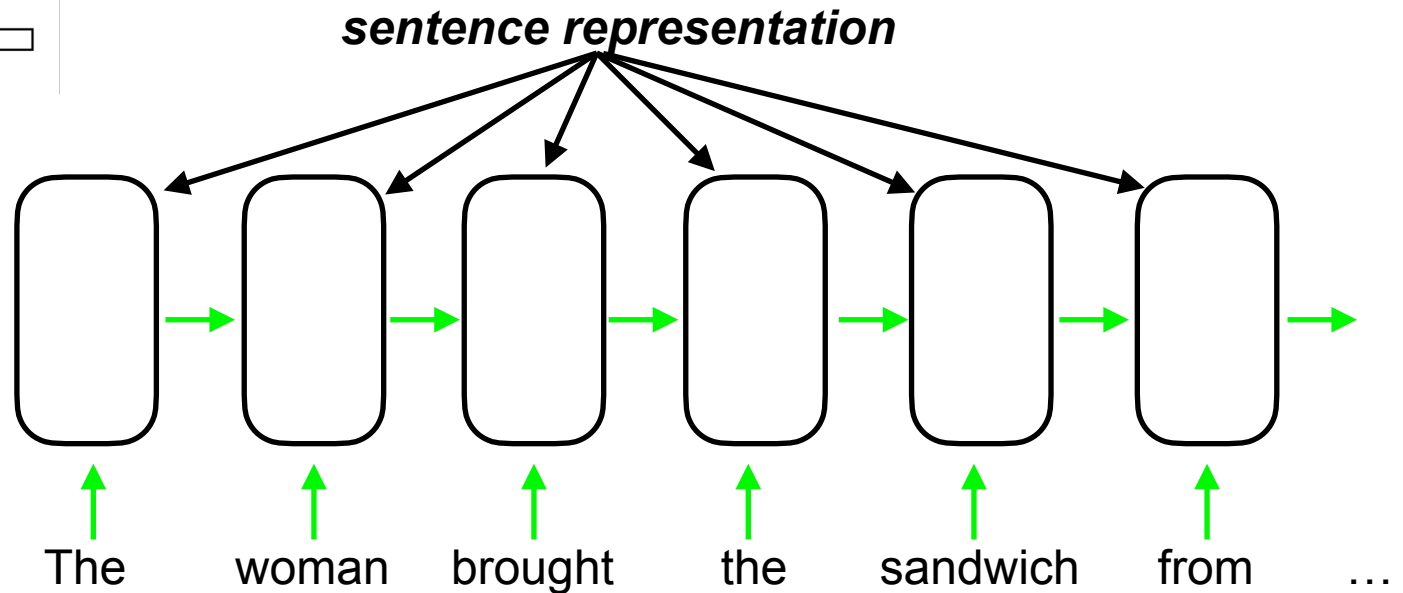

SRN "rolled up" and unrolled

- A “rolled-up” representation (Elman, 1990); and unrolled:



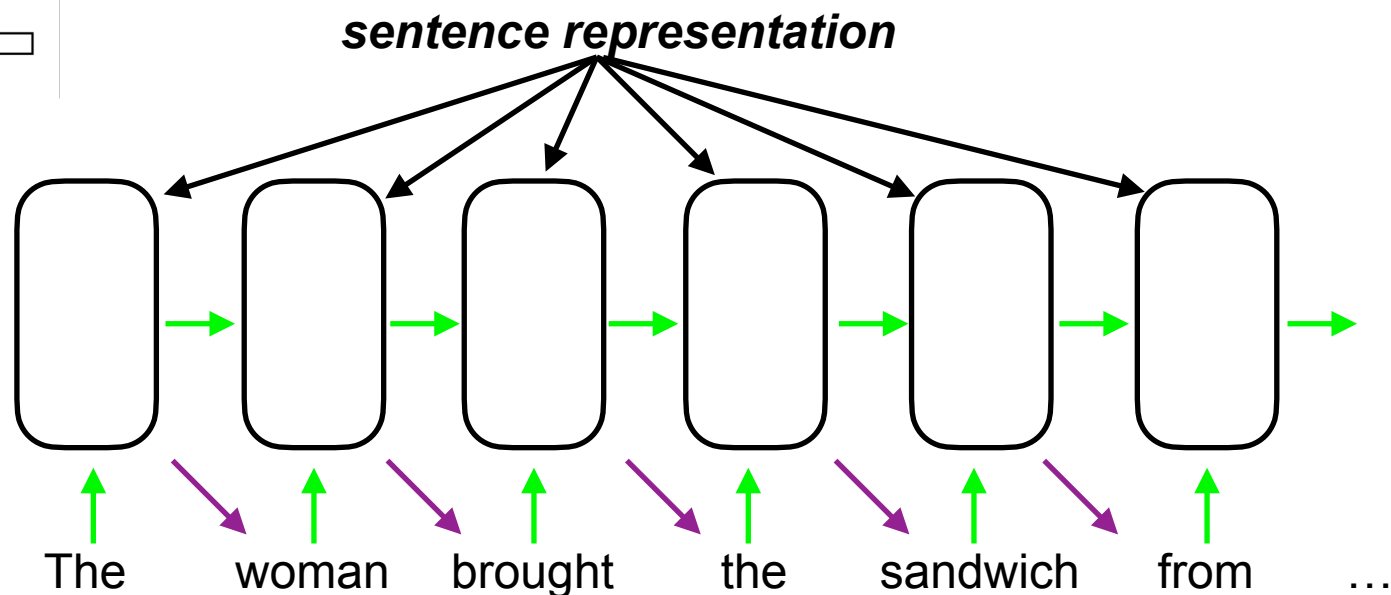

SRN "rolled up" and unrolled

- A “rolled-up” representation (Elman, 1990); and unrolled:



SRN "rolled up" and unrolled

- A “rolled-up” representation (Elman, 1990); and unrolled:



Predict!

Learning with artificial language input

TABLE 3
Categories of Lexical Items Used in Sentence Simulation

Category	Examples
NOUN-HUM	man, woman
NOUN-ANIM	cat, mouse
NOUN-INANIM	book, rock
NOUN-AGRESS	dragon, monster
NOUN-FRAG	glass, plate
NOUN-FOOD	cookie, break
VERB-INTRAN	think, sleep
VERB-TRAN	see, chase
VERB-AGPAT	move, break
VERB-PERCEPT	smell, see
VERB-DESTROY	break, smash
VERB-EAT	eat

TABLE 4
Templates for Sentence Generator

WORD 1	WORD 2	WORD 3
NOUN-HUM	VERB-EAT	NOUN-FOOD
NOUN-HUM	VERB-PERCEPT	NOUN-INANIM
NOUN-HUM	VERB-DESTROY	NOUN-FRAG
NOUN-HUM	VERB-INTRAN	
NOUN-HUM	VERB-TRAN	NOUN-HUM
NOUN-HUM	VERB-AGPAT	NOUN-INANIM
NOUN-HUM	VERB-AGPAT	
NOUN-ANIM	VERB-EAT	NOUN-FOOD
NOUN-ANIM	VERB-TRAN	NOUN-ANIM
NOUN-ANIM	VERB-AGPAT	NOUN-INANIM
NOUN-ANIM	VERB-AGPAT	
NOUN-INANIM	VERB-AGPAT	
NOUN-AGRESS	VERB-DESTROY	NOUN-FRAG
NOUN-AGRESS	VERB-EAT	NOUN-HUM
NOUN-AGRESS	VERB-EAT	NOUN-ANIM
NOUN-AGRESS	VERB-EAT	NOUN-FOOD







(Elman, 1990)

Used *localist* word representations

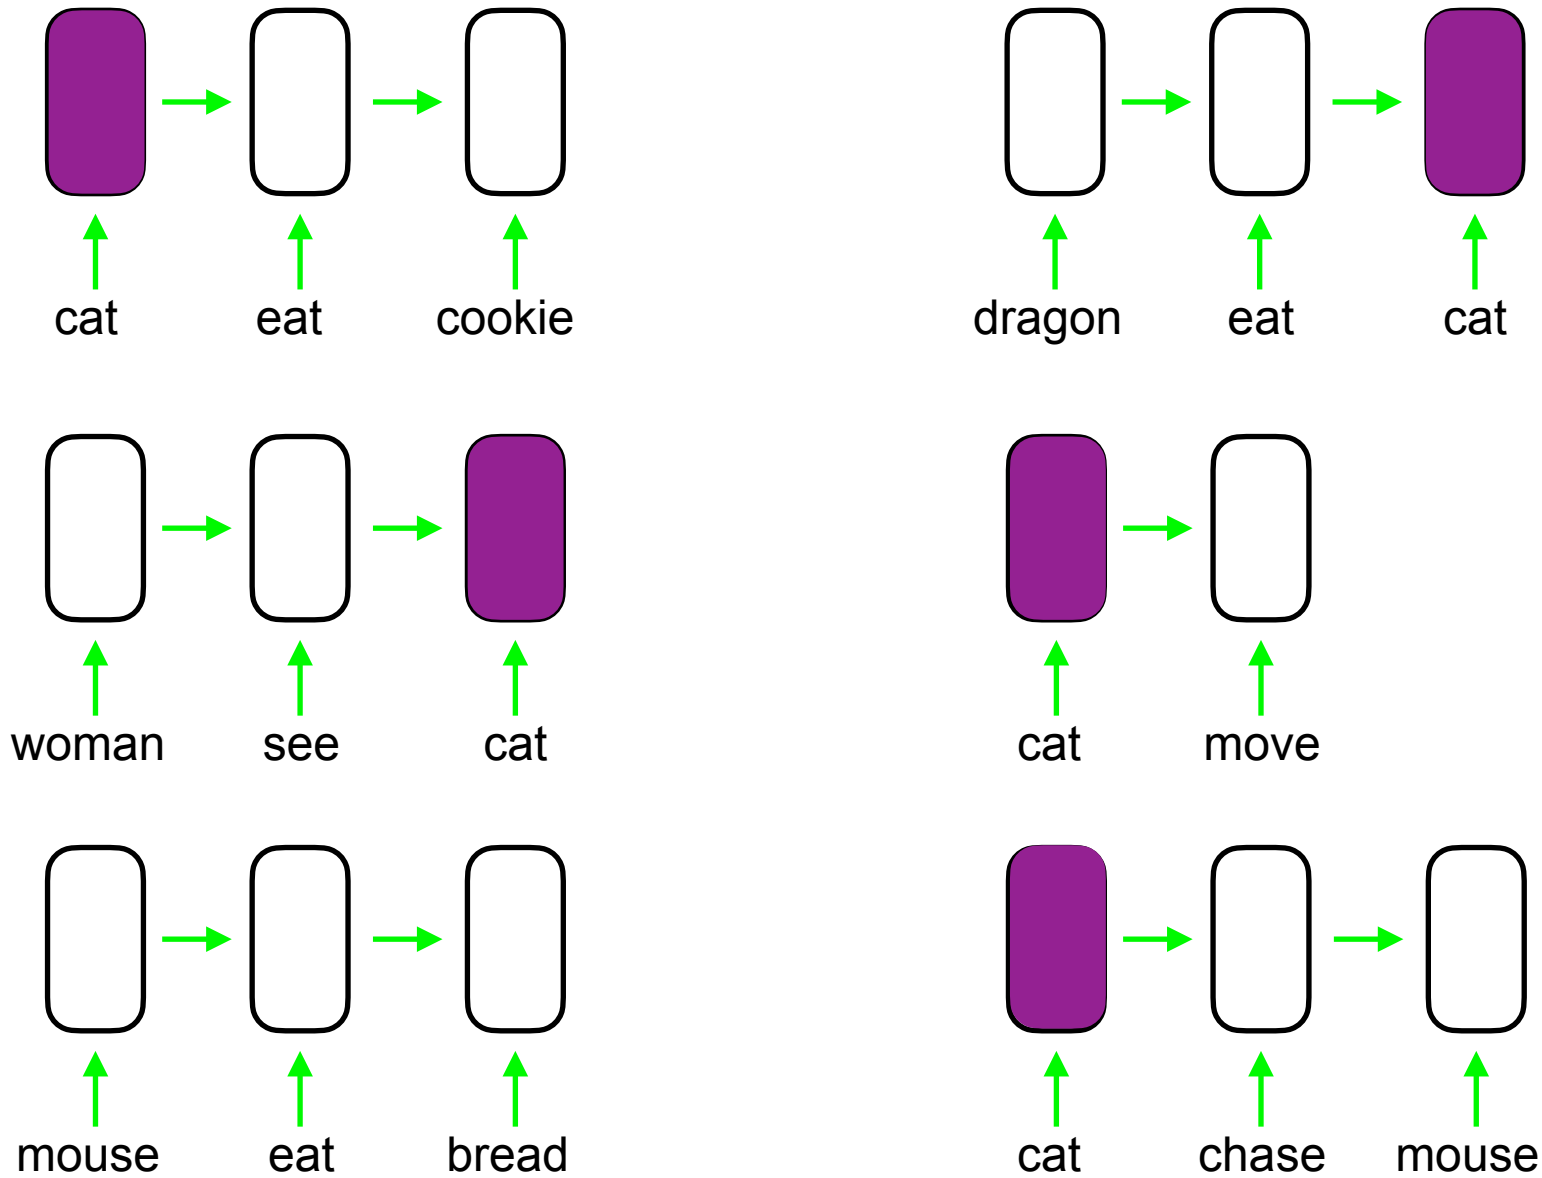
Fragment of Training Sequences for Sentence Simulation

Input	Output
00000000000000000000000000000010 (woman)	00000000000000000000000000000010000 (smash)
00000000000000000000000000000010000 (smash)	0000000000000000000000000000001000000000 (plate)
0000000000000000000000000000001000000000 (plate)	000001000000000000000000000000000000 (cat)
000001000000000000000000000000000000 (cat)	000000000000000000000000000000100000000000 (move)
000000000000000000000000000000100000000000 (move)	0000000000000000000000000000001000000000000 (man)
0000000000000000000000000000001000000000000 (man)	00010000000000000000000000000000000000 (break)
00010000000000000000000000000000000000 (break)	00001000000000000000000000000000000000 (car)
00001000000000000000000000000000000000 (car)	01000000000000000000000000000000000000 (boy)
01000000000000000000000000000000000000 (boy)	000000000000000000000000000000100000000000 (move)
000000000000000000000000000000100000000000 (move)	000000000000000000000000000000100000000000 (girl)
000000000000000000000000000000100000000000 (girl)	0000000000000000000000000000001000000000000 (eat)
000000000000000000000000000000100000000000 (eat)	00100000000000000000000000000000000000 (bread)
00100000000000000000000000000000000000 (bread)	000000000000000000000000000000100000000000 (dog)
000000000000000000000000000000100000000000 (dog)	000000000000000000000000000000100000000000 (move)
000000000000000000000000000000100000000000 (move)	000000000000000000000000000000100000000000 (mouse)
000000000000000000000000000000100000000000 (mouse)	000000000000000000000000000000100000000000 (mouse)
000000000000000000000000000000100000000000 (mouse)	000000000000000000000000000000100000000000 (move)
000000000000000000000000000000100000000000 (move)	10000000000000000000000000000000000000 (book)
10000000000000000000000000000000000000 (book)	0000000000000000000000000000001000000000000 (lion)

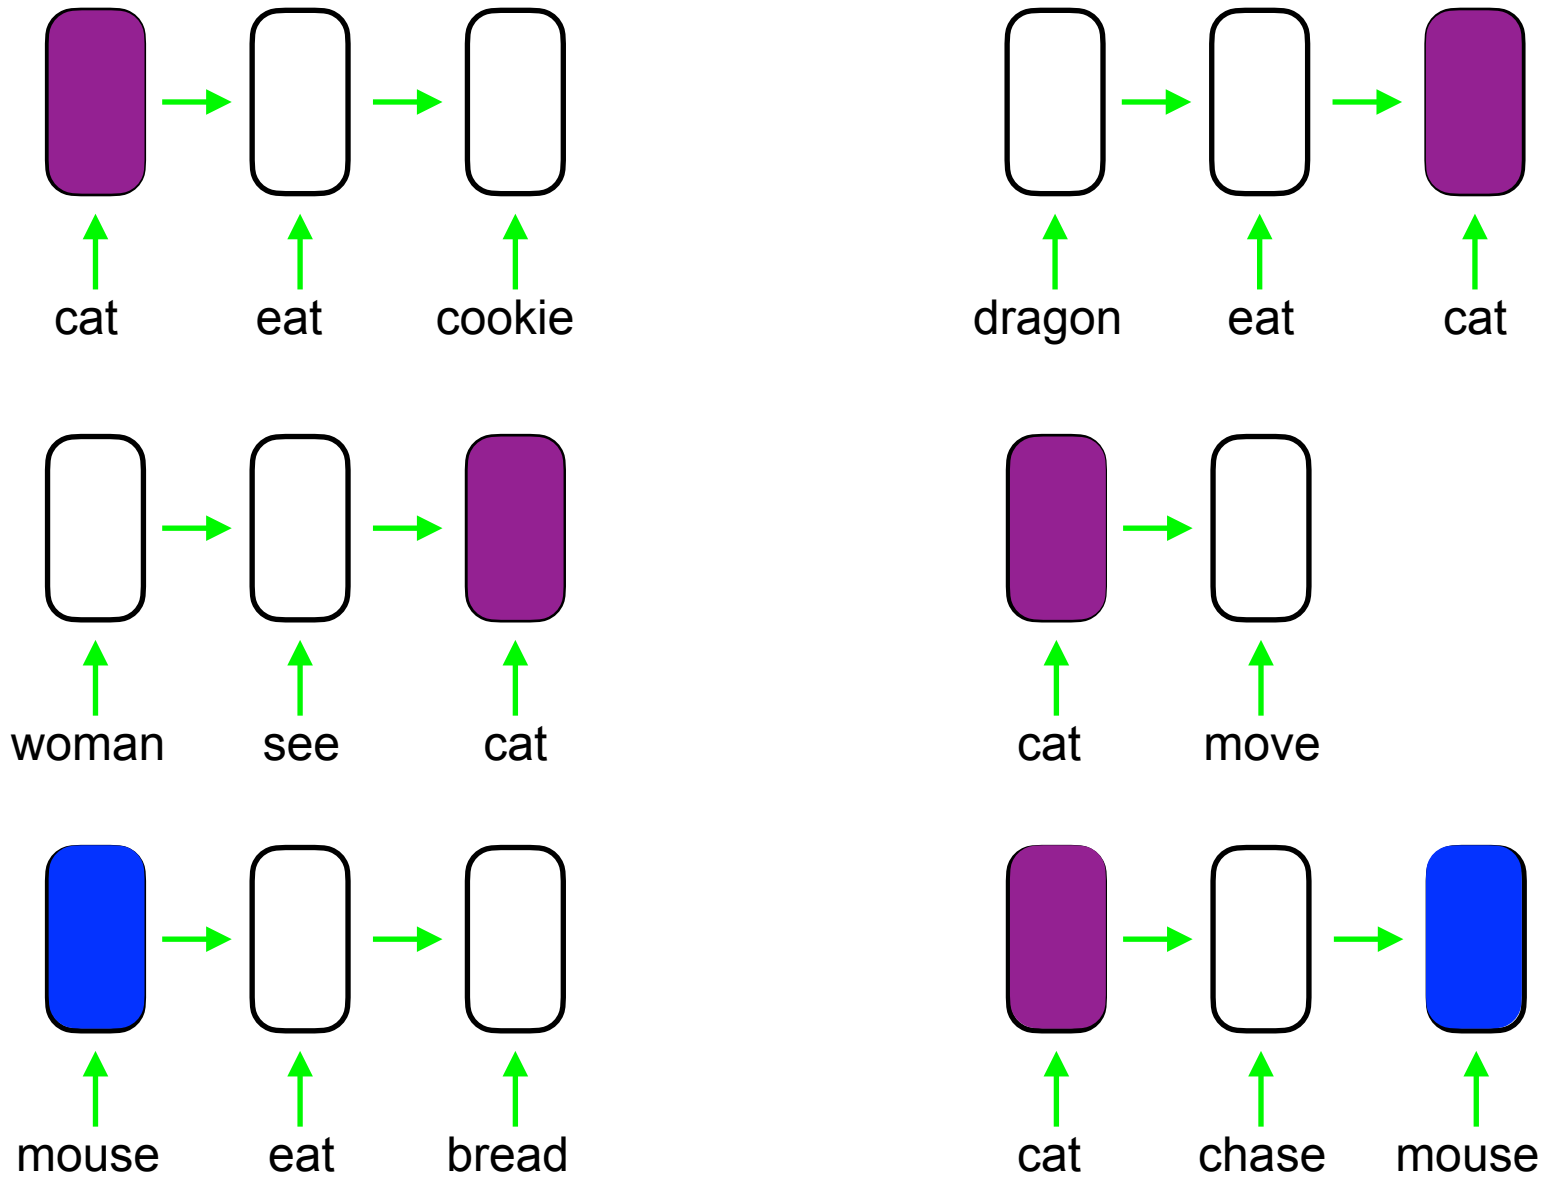
Learning word classes




Learning word classes



Learning word classes



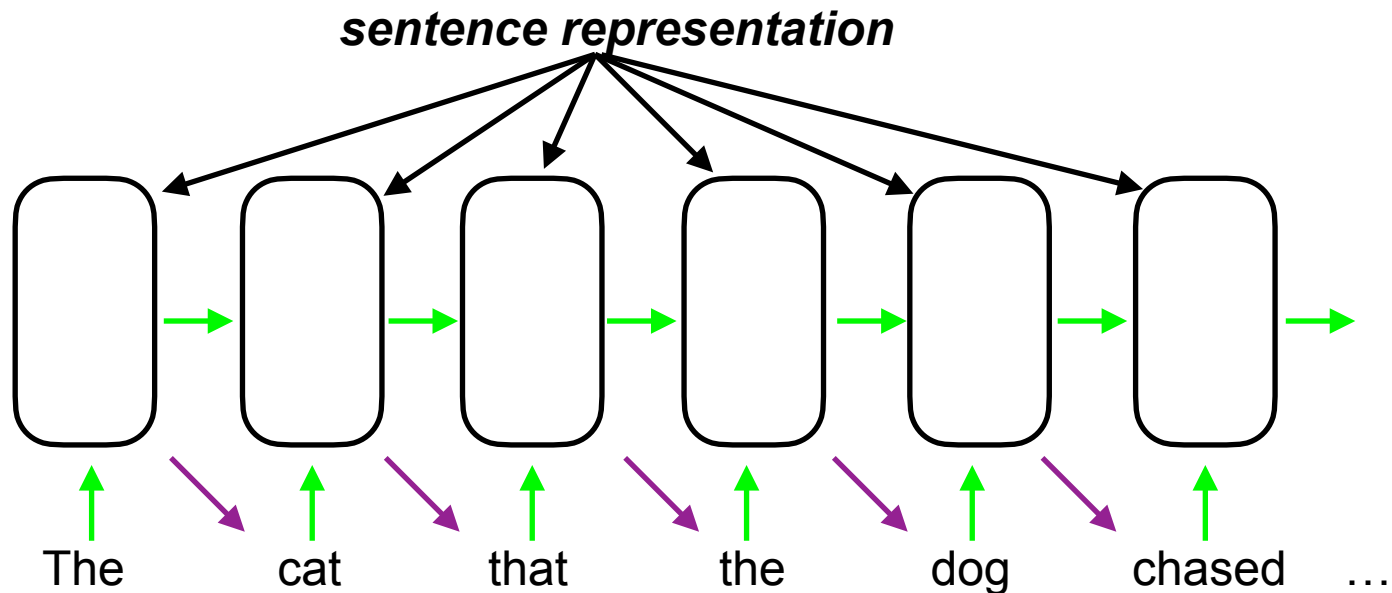
Discovered similarity structure of words




(Elman, 1990)

Beyond the simple recurrent network

- The SRN has a very strong ***linear locality bias***
- But natural language syntax is characterized by ***hierarchical structure***
- SRNs can learn hierarchy (Elman, 1991), but ***it is hard***—their inductive bias disfavors it




More sophisticated recurrent units




More sophisticated recurrent units

- Another view of an unrolled SRN:




More sophisticated recurrent units

- Another view of an unrolled SRN:




More sophisticated recurrent units

- Another view of an unrolled SRN:

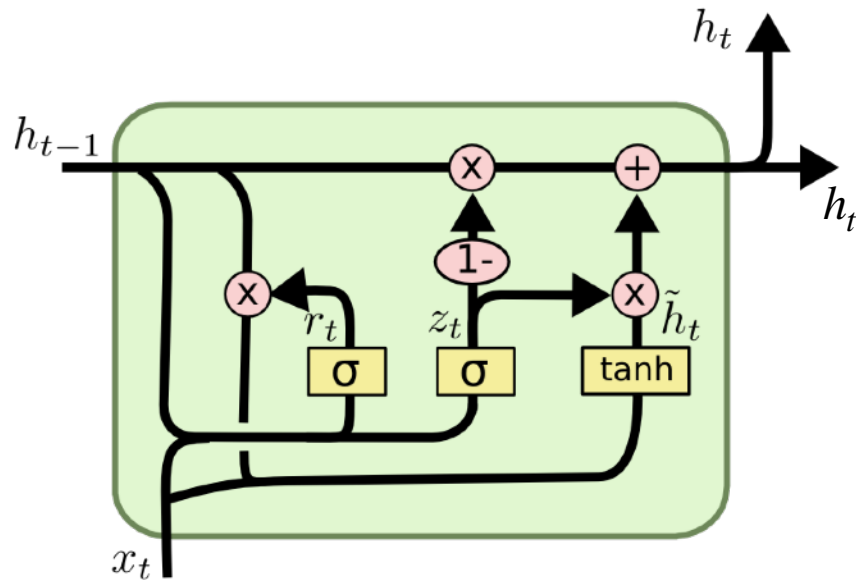


- Keep the recurrent structure and “swap in” a new unit:




(LSTM; Hochreiter & Schmidhuber, 1997)

Gated Recurrent Unit (GRU) architecture



Gated Recurrent Unit (GRU) architecture



$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

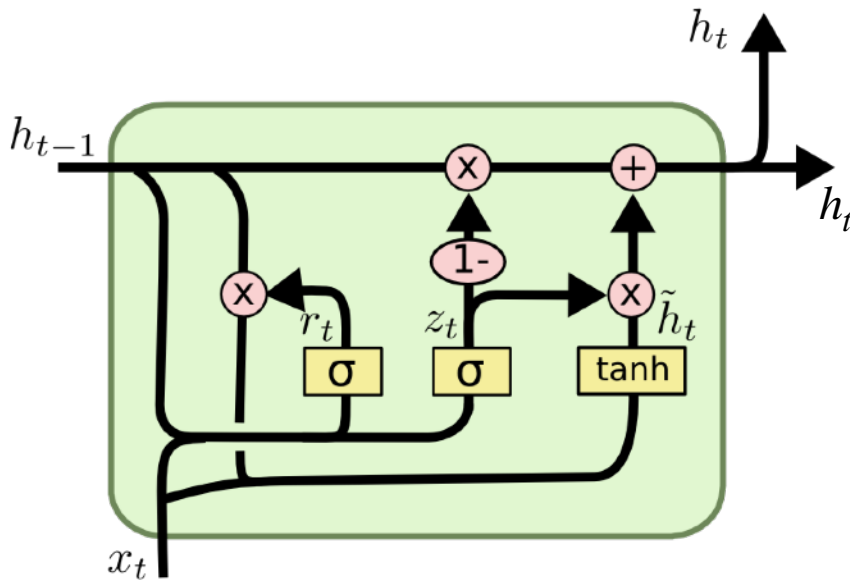
$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

Gated Recurrent Unit (GRU) architecture

logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$



$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$


$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

Gated Recurrent Unit (GRU) architecture

logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$



$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

Gated Recurrent Unit (GRU) architecture

logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$


$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$

element-wise multiplication

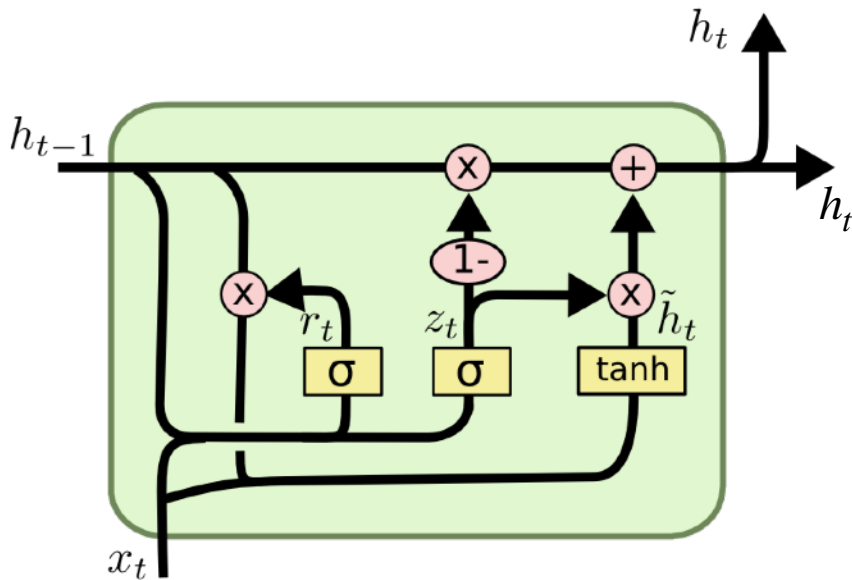
$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$(\text{e.g., } \langle 1, 2, 3 \rangle \odot \langle 0.5, 2, 1 \rangle = \langle 0.5, 4, 3 \rangle)$$



Gated Recurrent Unit (GRU) architecture



logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$


element-wise multiplication

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$(\text{e.g., } \langle 1, 2, 3 \rangle \odot \langle 0.5, 2, 1 \rangle = \langle 0.5, 4, 3 \rangle)$$

Gated Recurrent Unit (GRU) architecture



logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$


element-wise multiplication

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$(\text{e.g., } \langle 1, 2, 3 \rangle \odot \langle 0.5, 2, 1 \rangle = \langle 0.5, 4, 3 \rangle)$$

Gated Recurrent Unit (GRU) architecture



logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$


element-wise multiplication

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$(\text{e.g., } \langle 1, 2, 3 \rangle \odot \langle 0.5, 2, 1 \rangle = \langle 0.5, 4, 3 \rangle)$$

Gated Recurrent Unit (GRU) architecture



logistic/sigmoid activation function

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1})$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W} \mathbf{x}_t + \mathbf{U}(\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$


element-wise multiplication

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1})$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$(\text{e.g., } \langle 1, 2, 3 \rangle \odot \langle 0.5, 2, 1 \rangle = \langle 0.5, 4, 3 \rangle)$$


Long short-term memory (LSTM) units



(Hochreiter & Schmidhuber, 1997)

Inside the LSTM unit

- The “hidden layer” \mathbf{h}_{t-1} was used to predict element t of the sequence
- It now gets passed through a “forget gate”




$$\mathbf{f}_t = \sigma(\mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{W}_f \mathbf{x}_t)$$

(Hochreiter & Schmidhuber, 1997)

visualization due to Christopher Olah, <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Inside the LSTM unit

- Other information from h_{t-1} gets put into the memory store




$$\mathbf{i}_t = \sigma(\mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{W}_i \mathbf{x}_t)$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{U}_C \mathbf{h}_{t-1} + \mathbf{W}_C \mathbf{x}_t)$$

Inside the LSTM unit


- That information gets integrated into the memory store (which also gets passed on to the future)



$$\mathbf{C}_t = \mathbf{f}_t \odot \mathbf{C}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{C}}_t$$

Inside the LSTM unit


- Finally, we determine the new hidden layer to predict input $t+1$



$$\mathbf{o}_t = \sigma(\mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{W}_o \mathbf{x}_t)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{C}_t)$$

The LSTM unit, complete



$$\mathbf{f}_t = \sigma(\mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{W}_f \mathbf{x}_t)$$

$$\mathbf{i}_t = \sigma(\mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{W}_i \mathbf{x}_t)$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{U}_C \mathbf{h}_{t-1} + \mathbf{W}_C \mathbf{x}_t)$$

$$\mathbf{C}_t = \mathbf{f}_t \odot \mathbf{C}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{C}}_t$$

$$\mathbf{o}_t = \sigma(\mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{W}_o \mathbf{x}_t)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{C}_t)$$

Learning the classic counting language

$$a^n b^n$$

Learning the classic counting language

$a^n b^n$

Easily generable with a context-free grammar:

$S \rightarrow a b$

$S \rightarrow a S b$

Learning the classic counting language

$a^n b^n$

Easily generable with a context-free grammar:

$\wedge ab\$$

$S \rightarrow a \ b$

$\wedge aabb\$$

$S \rightarrow a \ S \ b$

$\wedge aaabbb\$$

$\wedge aaaabbbb\$$

$\wedge aaaaaabbbbbb\$$

$\wedge aaaaaaabbbbbbb\$$

$\wedge aaaaaaaaaabbbbbbbbbb\$$

\vdots

$\wedge aaaaaaaaaaaaaaaaaaaaaaaaaabbbbbbbbbbbbbbbbbbbbbbbbbb\$$

Learning the classic counting language

$a^n b^n$

Easily generable with a context-free grammar:

$S \rightarrow a \ b$
 $S \rightarrow a \ S \ b$

$^ab\$$

$^aabb\$$

$^aaabbb\$$

$^aaaabbbb\$$

$^aaaaabbbbb\$$

$^aaaaaabbbbbbb\$$

$^aaaaaaaaabbbbbbbbbb\$$

\vdots

$^aaaaaaaaaaaaaaaaaaaaaaaaabbbbbbbbbbbbbbbbbbbbbbbbbb\$$

D_{train}

Learning the classic counting language

$a^n b^n$

Easily generable with a context-free grammar:

$S \rightarrow a \ b$
 $S \rightarrow a \ S \ b$

$^{\wedge}ab\$$

$^{\wedge}aabb\$$

$^{\wedge}aaabbb\$$

$^{\wedge}aaaabbbb\$$

$^{\wedge}aaaaabbbbb\$$

$^{\wedge}aaaaaabbbbbb\$$

$^{\wedge}aaaaaaaaabbbbbbbbbb\$$


\vdots

$^{\wedge}\underbrace{aaaaaaaaaaaaaaaaaaaaaaaaa}_{N=20}\underbrace{bbbbbbbbbbbbbbbbbbbbbbbbbb}_{N=20}\$$

$N=20$


$N=20$

Training recurrent architectures on $a^n b^n$




Hidden & cell state contents


SRN hidden state




LSTM hidden state



GRU hidden state



LSTM cell state



Summary

- Mechanisms for neural networks at the sentence level:
 - Learned word embeddings
 - Recurrent state representation
- Different units used for recurrent state representation:
 - Simple recurrent network (SRN)
 - Gated recurrent unit (GRU)
 - Long short-term memory (LSTM)
- For classic counting language, LSTM works the best