


Noisy-channel sentence comprehension theory II




Roger Levy
9.19: Computational Psycholinguistics
22 November 2021

Today's agenda

- Explaining patterns of global utterance interpretation in the noisy-channel sentence processing theory
- Explaining "structural forgetting" effects by combining the noisy-channel theory with surprisal

Noisy-channel theory of language processing



Simple question-answering

The woman lost the diamond.

Did the woman lose something?

Yes

The ball kicked the girl.

Did the girl kick something?

No

The businessman benefited from the tax law.

Did the tax law benefit from anything?

No

The cook baked a cake Lucy.

Was something baked for Lucy?

No

(Yes?)


Over 2/3 of answers!

Noisy-channel semantic interpretation?

I ← The cook baked a cake Lucy.

m? Was something baked for Lucy?

Information



In two semantically plausible "neighbor" sentences, the answer is "yes":

The cook baked a cake Lucy.


Hypothesized noise operation: deletion

for

The cook baked a cake Lucy.

Hypothesized noise operation: exchange Lucy a cake

Predictions for implausible sentences



$$P(m | I) \propto P(I | m)P(m)$$

Noise operation Plausibility

Non-literal interpretation?

Implausible

Double Object/Benefactive-for alternation

The cook baked a cake Lucy.

Yes Yes

The cook baked Lucy for a cake.

Yes Yes

Plausible

The cook baked Lucy a cake.

No No

The cook baked a cake for Lucy.

No No

*Deletion/
insertion Exchange*

Implausible

Active/Passive alternation

The ball kicked the girl.

No Yes

The girl was kicked by the ball.

No Yes

Plausible

The girl kicked the ball.

No Yes

The ball was kicked by the girl.

No Yes

Literal vs. non-literal interpretation rates

Non-literal interpretations for implausible sentences?

Insertion/Deletion

Yes

Exchange

Yes



Insertion/Deletion

No

Exchange

Yes

Proportion of literal responses



Five alternations in an insertion/deletion model

English constructions	Change	Implausible version
1. Active/passive	Two insertions Two deletions	c. The girl <u>was</u> kicked <u>by</u> the ball. (passive) d. The ball kicked the girl. (active)
2. Subject-locative/ object-locative	One deletion, one insertion One insertion, one deletion	c. The table jumped <u>onto</u> a cat. (object-locative) d. <u>Onto</u> the cat jumped a table. (subject-locative)
3. Transitive/intransitive	One insertion One deletion	c. The tax law benefited <u>from</u> the businessman. (intransitive) d. The businessman benefited the tax law. (transitive)
4. DO/PO goal	One insertion One deletion	c. The mother gave the daughter <u>to</u> the candle. (PO-goal) d. The mother gave the candle the daughter. (DO-goal)
5. DO/PO benefactive	One insertion One deletion	c. The cook baked Lucy <u>for</u> a cake. (PO-benef) d. The cook baked a cake Lucy. (DO-benef)

c=inferred insertion d=inferred deletion


Five alternations in an insertion/deletion model

$$P(m|I) \propto P(I|m)P(m)$$

Noise operation Plausibility

Base experiment

20 experimental items,
60 plausible &
grammatically normal
fillers → 10/80
implausible trials



Fillers with syntactic errors

"A legislator lied to the
consultant a new bill"

"A bystander was the
fireman by rescued in
the nick of time"

Many implausible trials

100 experimental
items, 60 plausible &
grammatically normal
fillers → 50/160
implausible trials

Inferring deletions versus insertions

$$P(m | I) \propto \frac{P(I | m)}{\text{Noise operation}} P(m)$$

Noise operation Plausibility



Delete



Choose deletion location

The cook baked a cake ~~for~~ Lucy.



The cook baked a cake Lucy.



Insert



Choose insertion location

The cook baked Lucy a cake.



Choose what to insert



for

The cook baked Lucy for a cake.

Noisy-channel prediction: inferring deletions should be intrinsically easier than inferring insertions!

Five alternations in an insertion/deletion model

$$P(m|I) \propto P(I|m)P(m)$$

Noise operation ↑ Plausibility ↑

Base experiment

20 experimental items,
60 plausible &
grammatically normal
fillers → 10/80
implausible trials


Fillers with syntactic errors

"A legislator lied to the
consultant a new bill"

"A bystander was the
fireman by rescued in
the nick of time"

Many implausible trials

100 experimental
items, 60 plausible &
grammatically normal
fillers → 50/160
implausible trials



In the real world (2008)



I'm not going to solely
blame all of man's activities
on changes in climate.

Sarah Palin (images credit Gage Skidmore)



I'm not one to attribute every activity
of man to climate change.

Corpora of speech errors

Anticipations

John dropped his cuff of coffee

reek long race

Perseverations

John gave the goy (=gave the boy)

Spanish speaping people

teep a cape (=keep a tape)

Exchanges

the nipper is zarrow

Fancy getting your model renosed (=nose remodeled)

Revisiting the possibility of exchanges

This is a problem that I need to talk about Joe with.

- An occasional speech error of mine that I've noticed for years, but that no one ever notices me make
- Extraordinarily unlikely under an insertions/deletions noise model
- But reasonably likely if word **exchanges** are admitted

The package fell from the table to the floor. [plausible; canonical]


The package fell to the floor from the table. [plausible; non-canonical]

The package fell from the floor to the table. [implausible; canonical]


The package fell to the table from the floor. [implausible; non-canonical]

Did something fall to the floor?

Exchanges in the noise model



Probing inferred intended utterances



Experiment "cover story": read transcriptions of speech that might have errors, retype with edits if they think the speaker might have meant something else

The ball kicked the girl. The judge gave the athlete to the prize.

The ball kicked the girl.
No error


The judge gave the athlete the prize.
Insertion

The girl kicked the ball.
Exchange

The judge gave the athlete a prize.
Insertion

The ball was kicked by the girl. The judge gave the prize to the athlete.
Deletion
Exchange

Probing inferred intended utterances



Noisy-channel interpretation summary

- The noisy-channel framework suggests investigating global interpretations as well as incremental processing
- "Non-literal" interpretations can be very frequent for the right stimuli
- Interpretations broadly follow Bayesian principle of trade-off between prior and likelihood
 - Deletions easier to infer than insertions
 - Higher grammatical error rate in environment→more non-literal inference
 - More implausible sentences in environment→less non-literal inference
- *However*, status of exchange errors in the noise model remains a mystery


Structural Forgetting and the Noisy Channel

1. The apartment that the maid who the cleaning service sent over was well-decorated.
2. The apartment that the maid who the cleaning service sent over cleaned was well-decorated.

Structural Forgetting

1. *The **apartment** that the **maid** who the **cleaning service sent over was well-decorated.** 

2. The **apartment** that the **maid** who the **cleaning service sent over cleaned was well-decorated.** 



Structural Forgetting

1. *The **apartment** that the **maid** who the **cleaning service sent over was well-decorated.** 

2. The **apartment** that the **maid** who the **cleaning service sent over cleaned was well-decorated.** 



- **Structural forgetting effect:** part of the sentence is forgotten by the time you get to the end (Gibson & Thomas, 1999; Frazier, 1985; Fodor, p.c.)
- The ungrammatical sentence seems better than the grammatical one.
 - A "**grammaticality illusion**": how could we define grammaticality in this case?

Gibson & Thomas 1999: whole-sentence reading

The ancient manuscript that the graduate student who the new card catalog had confused a great deal was studying in the library was missing a page.

Christiansen & MacDonald 2009: word-by-word self-paced reading, follows by rating

The chef who the waiter who the busboy offended appreciated admired the musicians.



Structural Forgetting

1. *Die Wohnung, die das Zinайд who the cleaning service
Reinigungswasser well decorated, 👍 war gut eingerichtet. 👎
2. Die Wohnung, die das Zinайд who the cleaning service
Reinigung cleaned was well decorated, 👎 war gut eingerichtet. 👍

- But the effect is **language-dependent** (Vasishth et al., 2010; Frank et al., 2016).
 - In German (and Dutch), people prefer 2 over 1.
 - What is the difference between English and German?
 - Frank et al. (2016) show that at recurrent neural network gives higher probability to (1) in English, but (2) in German.
 - But why?

Structural Forgetting

1. *The **apartment** that the **maid** who the **cleaning service sent over was well-decorated.** 
 2. The **apartment** that the **maid** who the **cleaning service sent over cleaned was well-decorated.** 
- These contexts are more common in German than English (Roland et al., 2007).
 - English: the maid [that cleaned the apartment] **80%**
the apartment [that the maid cleaned] **20%**
 - German: das Dienstmädchen, [das die Wohnung reinigte] die Wohnung, [die das Dienstmädchen reinigte]

Noisy-Context Surprisal Account of Structural Forgetting

- Structural forgetting means the ungrammatical sentence with two verbs is **easier to process** than the grammatical sentence with three verbs:

C(The apartment that the maid who the cleaning service
NOUN THAT NOUN THAT NOUN VERB VERB) <
sent over was well-decorated.) <
C(NOUN(2 VERBS) NOUN THAT VERB VERB VERB VERB)
C(The apartment that the maid who the cleaning service
sent over cleaned was well-decorated.)

Noisy-Context Surprisal Account of Structural Forgetting

$$C(2 \text{ VERBS}) < C(3 \text{ VERBS})$$



- Correct noise based on prior about the language.
- Higher probability for verb-final RCs in German,
 - so more likely to make the right prediction.

Noisy-Context Surprisal Account of Structural Forgetting


- Futrell & Levy (2017) demonstrate that this works for toy grammars of English and German.

Rule	Probability						
$S \rightarrow NP\ VERB$	1		NOUN	VERB			
$NP \rightarrow NOUN$	$1-m$		NOUN	PREP	NOUN	VERB	
$NP \rightarrow NOUN\ RC$	mr		NOUN	THAT	VERB	NOUN	VERB
$NP \rightarrow NOUN\ PP$	$m(1-r)$		NOUN	THAT	NOUN	VERB	VERB
$PP \rightarrow PREP\ NP$	1		NOUN	THAT	NOUN	THAT	NOUN...
$RC \rightarrow THAT\ VERB\ NP$	s						
$RC \rightarrow THAT\ NP\ VERB$	$1-s$						



Plus **deletion noise**: every token in the context is forgotten (deleted) with probability d

Noisy-Context Surprisal Account of Structural Forgetting

- Setting the verb-final RC rate to 100% for German and 20% for English (Roland et al., 2007),
- we find surprisal differences matching the forgetting effect:



Noisy-Context Surprisal Account of Structural Forgetting



Vasishth et al. (2010)

Robustness to choice of model parameters


m Modifier probability

s Probability of English RC being verb-final

d Probability of context token deletion



= English+German-like pattern



Noisy-Context Surprisal Account of Structural Forgetting

- Probability that a context is remembered depends on its prior probability.
 - Noisy-context surprisal *explains* the behavior of the RNN in Frank et al. (2016): the RNN is using a lossily compressed / noisy representation of context.
- The model has an explicit grammar (competence), but cannot apply it correctly (performance).

Summary for today

- Applied noisy-channel sentence processing theory to global utterance interpretation
 - Broad support for noisy-channel predictions in patterns of "non-literal" utterance interpretation
 - A puzzle remains regarding the role of *exchange errors* in comprehenders' noise model
- Examined cross-linguistic pattern of "structural forgetting" effects, offered an account by combining noisy-channel theory + surprisal theory

References

- Christiansen, M. H., & MacDonald, M. C. (2009). A usage-based approach to recursion in sentence processing. *Language Learning*, 51(Suppl. 1), 126–161.
- Ferreira, F. (2003). The misinterpretation of noncanonical sentences. *Cognitive Psychology*, 47(2), 164–203.
- Frank, S. L., Trompenaars, T., & Vasishth, S. (2016). Cross-linguistic differences in processing double-embedded relative clauses: Working-memory constraints or language statistics? *Cognitive Science*, 40(3), 554–578.
- Frazier, L. (1985). Syntactic complexity. *Natural language parsing: Psychological, computational, and theoretical perspectives*, 129–189.
- Fromkin, V. A. (1971). The non-anomalous nature of anomalous utterances. *Language*, 51, 696–719.
- Futrell, R., Gibson, E., & Levy, R. P. (2020). Lossy-context surprisal: An information-theoretic model of memory effects in sentence processing. *Cognitive Science*, 44, 1–54.
- Futrell, R., & Levy, R. (2017). Noisy-context surprisal as a human sentence processing cost model. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics (EACL)* (pp. 688–698).
- Garrett, M. F. (1975). The analysis of sentence production. *Psychology of Learning and Motivation*, 9, 133–177.
- Gibson, E., Bergen, L., & Piantadosi, S. T. (2013). Rational integration of noisy evidence and prior semantic expectations in sentence interpretation. *Proceedings of the National Academy of Sciences*, 110(20), 8051–8056.
- Gibson, E., & Thomas, J. (1999). The perception of complex ungrammatical sentences as grammatical. *Language & Cognitive Processes*, 14(3), 225–248.
- Levy, R. (2008). A noisy-channel model of rational human sentence comprehension under uncertain input. In *Proceedings of the 13th conference on Empirical Methods in Natural Language Processing* (pp. 234–243). Waikiki, Honolulu.
- Poppels, T., & Levy, R. (2016). Structure-sensitive noise inference: Comprehenders expect exchange errors. In *Proceedings of the 38th annual meeting of the Cognitive Science Society* (pp. 378–383).
- Roland, D., Dick, F., & Elman, J. L. (2007). Frequency of basic English grammatical structures: A corpus analysis. *Journal of Memory and Language*, 57, 348–379.
- Ryskin, R., Futrell, R., Kiran, S., & Gibson, E. (2018). Comprehenders model the nature of noise in the environment. *Cognition*, 181, 141–150.
- Shannon, C. (1948). A mathematical theory of communication. *Bell Systems Technical Journal*, 27(4), 623–656.
- Vasishth, S., Suckow, K., Lewis, R. L., & Kern, S. (2010). Short-term forgetting in sentence comprehension: Crosslinguistic evidence from verb-final structures. *Language & Cognitive Processes*, 25(4), 533–567.

Prediction 2: hallucinated garden paths

- Try reading the sentence below:

While the clouds crackled, above the glider soared a magnificent eagle.

- There's a garden-path clause in this sentence...
- ...but it's interrupted by a comma.
- Readers are ordinarily very good at using commas to guide syntactic analysis:

While the man hunted, the deer ran into the woods

While Mary was mending the sock fell off her lap


- “With a comma after *mending* there would be no syntactic garden path left to be studied.” (Fodor, 2002)
- We'll see that the story is slightly more complicated.

Prediction 2: hallucinated garden paths

While the clouds crackled, above the glider soared a magnificent eagle.



- This sentence is comprised of an initial intransitive subordinate clause...
- ...and then a main clause with *locative inversion*.
(c.f. a magnificent eagle soared above the glider)
- Crucially, the main clause's initial PP would make a great dependent of the subordinate verb...
- ...but doing that *would require the comma to be ignored*.
- Inferences through ...*glider* should thus involve a tradeoff between perceptual input and prior expectations



- Inferences as probabilistic paths through the sentence:
 - Perceptual cost of ignoring the comma
 - Unlikeliness of main-clause continuation after comma
 - Likeliness of postverbal continuation without comma
- These inferences together make *soared* very surprising!

$$P(w_i|\text{Context}) = \sum_{\text{Path}} P(w_i|\text{Path, Context})P(\text{Path}|\text{Context})$$

Prediction 2: hallucinated garden paths

- Two properties come together to create “hallucinated garden path”
 1. Subordinate clause into which the main-clause inverted phrase would fit well
 2. Main clause with locative inversion
- Experimental design: cross (1) and (2)

While the clouds crackled, above the glider soared a magnificent eagle.

While the clouds crackled, the glider soared above a magnificent eagle.

While the clouds crackled in the distance, above the glider soared a magnificent eagle.

While the clouds crackled in the distance, the glider soared above a magnificent eagle.

- The phrase *in the distance* fulfills a similar thematic role as above the glider for crackled
- Should reduce hallucinated garden-path effect

Prediction 2: Hallucinated garden paths

- Methodology: word-by-word self-paced reading

white-the-~~the~~-heads-crackled,-above-the-glider-soared-----

- Readers aren't allowed to backtrack
- So the comma is visually *gone* by the time the inverted main clause appears
- Simple test of whether beliefs about previous input can be revised


Model predictions

While the clouds
crackled, **above** the
glider soared a
magnificent eagle.

While the clouds **crackled**
in the distance, **above**
the glider soared a
magnificent eagle.


While the clouds
crackled, the glider
soared **above** a
magnificent eagle.

While the clouds **crackled**
in the distance, the
glider soared **above** a
magnificent eagle.



Results: whole sentence reading times

Processing boggle occurs exactly where predicted





Hallucinated garden-path summary

- The *at/toward* study showed that comprehenders *note the possibility of alternative strings and act on it*
- This study showed that comprehenders can actually *devote resources to grammatical analyses inconsistent with the surface string*

Hallucinated garden paths cont'd

- Sure, but punctuation's weird stuff
- What about *real words*?

I know that the desert trains could resupply the camp.



- At least sometimes, bias against N N interpretation

Hallucinated GPs with words

- We use a contextual bias against NN and toward NV to test for GP hallucinations involving wordform change

Could be “intern chauffeured”

The intern chauffeur for the governor hoped for more interesting work.
[NN, “dense” neighborhood]

The intern chauffeured for the governor but hoped for more interesting work.
[NV, “dense” neighborhood]


Could NOT be “inexperienced chauffeured”

The inexperienced chauffeur for the governor hoped for more interesting work.
[NN, “sparse” neighborhood]

Some interns chauffeured for the governor but hoped for more interesting work.
[NV, “sparse” neighborhood]

Results

- RT spike at disambiguating region for NN Dense



Structure of the noise model

- Gibson et al. (2013) explored noise model with restricted noise operations: word *insertions* and *deletions*

Sentence	Plausibility	Insertions	Deletions
The cook baked Lucy a cake.	Plausible	0	1
The cook baked Lucy <u>for</u> a cake.	Implausible	1	0
The cook baked a cake <u>for</u> Lucy.	Plausible	1	0
The cook baked a cake Lucy.	Implausible	0	1

Structure of the noise model

Consistently given literal interpretation

Sentence

The girl was kicked by the ball.

The ball kicked the girl.

The tax law benefited from the businessman.

The businessman benefited the tax law.

The cook baked Lucy for a cake.

The cook baked a cake Lucy.

Construction

passive

active

intransitive

transitive

Prepositional Object
(PO) benefactive

Double Object (DO)
benefactive

Edits

2I

2D

1I

1D


1I

1D


Often “corrected” to plausible interpretation inconsistent with literal meaning

Dependency length and noisy-channel surprisal


- Syntactic dependencies vary in linear distance



- Idea with long history: short dependencies preferred



Dependency lengths are short across languages!




Dependency lengths and the noisy channel

- Here: dependency length minimization can be derived from a combination of surprisal & noisy-channel theory

From noisy-channel & surprisal to dependency length minimization

noisy context



- Suppose we have an **increasing noise rate** the longer a word has been in memory.
- When "threw" is far from "out", then it is less likely to reduce the surprisal of "out": more likely to be affected by noise.
- Noisy-context surprisal increases when **words that predict each other are far apart**.
- We call this **information locality** (following Gildea & Jaeger, 2015).

Derivation of Information Locality


- Erasure noise decreases the influence of context:

$$C(w|\text{context}) \approx h(w) - \sum_{w' \in \text{context}} P(w' \text{ not erased}) \text{pmi}(w; w')$$

John threw the trash out

$h(\text{out}) - P(\text{John not erased}) \text{pmi}(\text{John}; \text{out})$

- $P(\text{threw not erased}) \text{pmi}(\text{threw}; \text{out})$
- $P(\text{the not erased}) \text{pmi}(\text{the}; \text{out})$
- $P(\text{trash not erased}) \text{pmi}(\text{trash}; \text{out})$



Derivation of Information Locality

- Noise decreases the influence of context:

$$C(w|\text{context}) \approx h(w) - \sum_{w' \in \text{context}} P(w' \text{ not erased}) \text{pmi}(w; w')$$

threw out

$h(\text{out}) - P(\text{threw not erased}) \text{ pmi}(\text{threw}; \text{out})$

```

0000011010101011001011010000111001111011100
010101011110101100101000000010101110010011
001001101010110011000101010010011010100110
1001001000010010100110111100100100100010001000
011000111001111000100111100100101101001010
11000010000110011000010101010011111111100
110100110011011110000001011000111001111010
010100010011111101101111100101001011000001
100111100100001011110001000110000111010001
001111010100111101110010100011100100100101
101011001000101110101000011110011101101101
11010100001100010011000101000100100100101000
001101100100001001001010100010100110000011
001001101001111011110100110100011010111000
1000001010101000101101001111010110101100111
011010100001100000110001000001111111111001
11010100001110100110111100001110001110010111
001110111100011101011110011011111100001110
0111100110011001101010110010111100100000000
011 P(threw not erased).pmi(threw;out)1110
110000010010011110110101001101011100001100
0011110101010110111111011110011001001010111
101010011110110110010000111100100100000000
111110111101001010000100101000101000001101
01010110010110100001110011110111001010101
111010110010100000001010111001001100100100110
10101100110001010100100110101001101010010010

```

Derivation of Information Locality

- Noise decreases the influence of context:

$$C(w|\text{context}) \approx h(w) - \sum_{w' \in \text{context}} P(w' \text{ not erased}) \text{pmi}(w; w')$$

threw **out**

$h(\text{out}) - P(\text{threw} \text{ not erased}) \text{ pmi}(\text{threw}; \text{out})$

- When context items are far, their cost-reducing influence decreases.
 - Similar to the concept of decay in cue effectiveness
(Qian & Jaeger, 2012)


00001101010110001011010000111001111011100
010101011110101100101000000010101110010011
0010011010101100110001010100100110100110
10010010000100101001101111001001000100001000
011000111001111000100111100100101101001010
11000010000110011000010101010011111111100
11010011001101110000001011000111001111010
010100010011111101101111100101001011000001
100111100100001011110001000110000111010001
001111010100111101110010100011100100100101
101011001000101110101000011110011101101101
110101000001100010011000101000100100101000
001101100100001001001010100010100110000011
001001101001111011110100110100011010111000
100001010101000101101001111010110101100111
011010100001100000110001000001111111111001
110101000011101001101110000111000111001011
001110111100011101011110011011111100001110
011110011001100110110101100101111001100000
011110010101111100110100011000000000111110
1100000100100111110101101001101011110001100
001111010101011011111101111001100010010111
101010011111010110010000111100100100000000
111110111101001010000100101000101000001101
010101010101010101010001010100010101010101
111010110010010000000101011001010101010101
01010110010001010100100110101010011010010010

C

Information Locality


- **Information locality:** prediction of processing difficulty when words that predict each other (have high mutual information) are far apart.
- How does this relate to **dependency locality?**
- Hypothesis: **Words in syntactic dependencies have high mutual information.**
 - If this is true, then we can see dependency locality effects as a subset of information locality effects.
- We will show that the hypothesis is true in dependency corpora.

Do Dependencies Have High Mutual Information?



- We calculated mutual information values over part-of-speech tags for pairs of words in the UD corpora.

Do Dependencies Have High Mutual Information?



Comprehension as exploration of input

- Broader ongoing goal: develop eye-movement control model integrating the insights discussed thus far:
 - Probabilistic linguistic knowledge
 - Uncertain input representations
 - Principles of adaptive, rational action
- *Reinforcement learning* is an attractive tool for this

A rational reader

- Very simple framework:
 - Start w/ prior expectations for text (linguistic knowledge)
 - Move eyes to get perceptual input
 - Update beliefs about text as visual arrives (Bayes' Rule)
- Add to that:
 - Set of *actions* the reader can take in discrete time
 - A *behavior policy*: how the model decides between actions

A first-cut behavior policy

- Actions: *keep fixating*; *move the eyes*; or *stop reading*
- Simple behavior policy with two parameters: α and β
- Define *confidence* in a character position as the probability of the most likely character

*From the closet, she pulled out a *acket for the upcoming game*


Confidence=0.59

$P(\text{jacket})=0.38$
 $P(\text{racket})=0.59$
 $P(\text{packet})=0.02$
...

- Move left to right, bringing up confidence in each character position until it reaches α
- If confidence in a previous character position drops below β , regress to it
- Finish reading when you're confident in everything

(Non)-regressive policies

- Non-regressive policies have $\beta=0$
- Hypothesis: non-regressive policies strictly dominated
- Test: estimate speed and accuracy of various policies on reading the the Schilling et al. (1998) corpus



Goal-based adaptation

- Open frontier: modeling the adaptation of eye movements to specific reader goals
- We set a *reward function*: relative value γ of speed (finish reading in T timesteps) versus accuracy (guess correct sentence with probability L)
- PEGASUS simplex-based optimization (Ng & Jordan, 2000)







γ	α	β
0.025		
0.1		
0.4		

- The method works, and gives intuitive results

Empirical match with human reading

- Benchmark measures in eye-movement modeling:

frequency



predicts size and
shape of all effects

predictability


Bicknell & Levy (2012)

Success at empirical benchmarks

- Other models (E-Z Reader, SWIFT) get these too, but *stipulate* rel'nship between word properties & “processing rate”
- We *derive* these relationships from simple principles of noisy-channel perception and rational action

Noisy-channel processing: summary

- Noisy-channel models help us understand
 - Basic capabilities of human language comprehension
 - Outstanding puzzles in syntactic processing
- These models open up a rich typology of new sentence processing effects
- There is growing evidence for these effects
- These models pose new theoretical opportunities and architectural challenges for the study of human linguistic cognition



References I

-  Bergen, L., Levy, R., & Gibson, E. (2012). Verb omission errors: Evidence of rational processing of noisy language inputs. In *Proceedings of the 34th annual meeting of the Cognitive Science Society* (pp. 1320–1325).
-  Bever, T. (1970). The cognitive basis for linguistic structures. In J. Hayes (Ed.), *Cognition and the development of language* (pp. 279–362). New York: John Wiley & Sons.
-  Bicknell, K., & Levy, R. (2010). A rational model of eye movement control in reading. In *Proceedings of the 48th annual meeting of the Association for Computational Linguistics* (pp. 1168–1178). Uppsala, Sweden.
-  Bicknell, K., & Levy, R. (2012a). Why long words take longer to read: The role of uncertainty about word length. In *Proceedings of the 3rd annual workshop on Cognitive Modeling and Computational Linguistics* (pp. 21–30).

References II

-  Bicknell, K., & Levy, R. (2012b). Word predictability and frequency effects in a rational model of reading. In *Proceedings of the 34th annual meeting of the Cognitive Science Society* (pp. 126–131). Sapporo, Japan.
-  Booth, T. L. (1969). Probabilistic representation of formal languages. In *IEEE conference record of the 1969 tenth annual symposium on switching and automata theory* (pp. 74–81).
-  Crocker, M., & Brants, T. (2000). Wide-coverage probabilistic sentence processing. *Journal of Psycholinguistic Research*, 29(6), 647–669.
-  Fodor, J. D. (2002). Psycholinguistics cannot escape prosody. In *Proceedings of the speech prosody conference*.

References IV

-  Futrell, R., & Levy, R. (2017). Noisy-context surprisal as a human sentence processing cost model. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics (EACL)* (pp. 688–698).
-  Futrell, R., Mahowald, K., & Gibson, E. (2015). Large-scale evidence of dependency length minimization in 37 languages. *Proceedings of the National Academy of Sciences*, 112(33), 10336–10341.
-  Gibson, E. (1998). Linguistic complexity: Locality of syntactic dependencies. *Cognition*, 68, 1–76.
-  Gibson, E. (2000). The dependency locality theory: A distance-based theory of linguistic complexity. In A. Marantz, Y. Miyashita, & W. O’Neil (Eds.), *Image, language, brain* (pp. 95–126). Cambridge, MA: MIT Press.

References V

-  Gibson, E., Bergen, L., & Piantadosi, S. T. (2013). Rational integration of noisy evidence and prior semantic expectations in sentence interpretation. *Proceedings of the National Academy of Sciences*, 110(20), 8051–8056.
-  Gibson, E., & Thomas, J. (1999). The perception of complex ungrammatical sentences as grammatical. *Language & Cognitive Processes*, 14(3), 225–248.
-  Gildea, D., & Jaeger, T. F. (2015). Human languages order information efficiently. *CoRR*, abs/1510.02823. arXiv: 1510.02823
-  Gildea, D., & Temperley, D. (2007). Optimizing grammars for minimum dependency length. In *Proceedings of the annual meeting of the association for computational linguistics*, Prague, Czech Republic.
-  Gildea, D., & Temperley, D. (2010). Do grammars minimize dependency length? *Cognitive Science*, 34, 286–310.

References VI

-  Hale, J. (2001). A probabilistic Earley parser as a psycholinguistic model. In *Proceedings of the second meeting of the north american chapter of the Association for Computational Linguistics* (pp. 159–166). Pittsburgh, Pennsylvania.
-  Hale, J. (2006). Uncertainty about the rest of the sentence. *Cognitive Science*, 30(4), 609–642.
-  Hawkins, J. A. (2004). *Efficiency and complexity in grammars*. Oxford University Press.
-  Itti, L., & Baldi, P. (2005). Bayesian surprise attracts human attention. In *Advances in Neural Information Processing Systems*.
-  Jelinek, F., & Lafferty, J. D. (1991). Computation of the probability of initial substring generation by stochastic context free grammars. *Computational Linguistics*, 17(3), 315–323.

References VII

-  Johnson, M. (1998). PCFG models of linguistic tree representations. *Computational Linguistics*, 24(4), 613–632.
-  Klein, D., & Manning, C. D. (2003). Accurate unlexicalized parsing. In *Proceedings of acl*.
-  Levy, R. (2008a). A noisy-channel model of rational human sentence comprehension under uncertain input. In *Proceedings of the 13th conference on Empirical Methods in Natural Language Processing* (pp. 234–243). Waikiki, Honolulu.
-  Levy, R. (2008b). Expectation-based syntactic comprehension. *Cognition*, 106(3), 1126–1177.

References VIII

-  Levy, R. (2011). Integrating surprisal and uncertain-input models in online sentence comprehension: Formal techniques and empirical results. In *Proceedings of the 49th annual meeting of the Association for Computational Linguistics: Human Language Technologies* (pp. 1055–1065).
-  Levy, R. (2013). Memory and surprisal in human sentence comprehension. In R. P. G. van Gompel (Ed.), *Sentence processing* (pp. 78–114). Hove: Psychology Press.
-  Levy, R., Bicknell, K., Slattery, T., & Rayner, K. (2009). Eye movement evidence that readers maintain and act on uncertainty about past linguistic input. *Proceedings of the National Academy of Sciences*, 106(50), 21086–21090.
-  MacDonald, M. C. (1993). The interaction of lexical and syntactic ambiguity. *Journal of Memory and Language*, 32, 692–715.

References IX

-  Mohri, M. (1997). Finite-state transducers in language and speech processing. *Computational Linguistics*, 23(2), 269–311.
-  Park, Y. A., & Levy, R. (2009). Minimal-length linearizations for mildly context-sensitive dependency trees. In *Proceedings of the 10th annual meeting of the North American chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT) conference* (pp. 335–343). Boulder, Colorado, USA.
-  Poppels, T., & Levy, R. (2016). Structure-sensitive noise inference: Comprehenders expect exchange errors. In *Proceedings of the 38th annual meeting of the Cognitive Science Society* (pp. 378–383).
-  Roland, D., Dick, F., & Elman, J. L. (2007). Frequency of basic English grammatical structures: A corpus analysis. *Journal of Memory and Language*, 57, 348–379.

References X

-  Smith, N. J., & Levy, R. (2013). The effect of word predictability on reading time is logarithmic. *Cognition*, 128(3), 302–319.
-  Staub, A. (2007). The parser doesn't ignore intransitivity, after all. *Journal of Experimental Psychology: Learning, Memory, & Cognition*, 33(3), 550–569.
-  Stolcke, A. (1995). An efficient probabilistic context-free parsing algorithm that computes prefix probabilities. *Computational Linguistics*, 21(2), 165–201.
-  Tabor, W., Galantucci, B., & Richardson, D. (2004). Effects of merely local syntactic coherence on sentence processing. *Journal of Memory and Language*, 50(4), 355–370.
-  Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science*, 268, 1632–1634.

References XI

- 
- Vasishth, S., Suckow, K., Lewis, R. L., & Kern, S. (2010).
Short-term forgetting in sentence comprehension:
Crosslinguistic evidence from verb-final structures.
Language & Cognitive Processes, 25(4), 533–567.